

personalized-medicine-redefining-cancer-treatmen

June 22, 2021

Cancer Treatment

```
[114]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import SGDClassifier
from imblearn.over_sampling import SMOTE
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
warnings.filterwarnings("ignore")

from mlxtend.classifier import StackingClassifier
```

```
from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
```

0.1 Exploring and reading the datafiles

```
[3]: # Loading training data
data_variants = pd.read_csv('data/training_variants')
# training_text dataset uses "/" as a separator and has the headers separated
↳ by ",", so we skip that row and declare headers
data_text = pd.read_csv("data/
↳ training_text", sep="\\|\\|", engine="python", names=["ID", "TEXT"], skiprows=1)
```

Exploring the datasets **data_variants**

```
[4]: data_variants.head()
```

```
[4]:
```

	ID	Gene	Variation	Class
0	0	FAM58A	Truncating Mutations	1
1	1	CBL	W802*	2
2	2	CBL	Q249E	2
3	3	CBL	N454D	3
4	4	CBL	L399V	4

Some data explanation

ID : row id used to link the mutation to the clinical evidence

Gene : the gene where this genetic mutation is located

Variation : the aminoacid change for this mutations

Class : class value 1-9, this genetic mutation has been classified on

This is be our result data

```
[5]: data_variants.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3321 entries, 0 to 3320
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ID          3321 non-null   int64
1   Gene        3321 non-null   object
2   Variation    3321 non-null   object
3   Class       3321 non-null   int64
dtypes: int64(2), object(2)
memory usage: 103.9+ KB
```

There are no null-values in the data_varuants

```
[6]: data_variants.describe()
```

```
[6]:
```

	ID	Class
count	3321.000000	3321.000000
mean	1660.000000	4.365854
std	958.834449	2.309781
min	0.000000	1.000000
25%	830.000000	2.000000
50%	1660.000000	4.000000
75%	2490.000000	7.000000
max	3320.000000	9.000000

Mathematical description of the dataset

```
[7]: # Checking dimention of data
# even though we already know from count and columns exploration (data_variants.
# head(3))
data_variants.shape
```

```
[7]: (3321, 4)
```

```
[8]: # Clecking column in above data set
# even though we already know from data_variants.head(3)
data_variants.columns
```

```
[8]: Index(['ID', 'Gene', 'Variation', 'Class'], dtype='object')
```

data_text

```
[9]: data_text.head(3)
```

```
[9]:
```

	ID	TEXT
0	0	Cyclin-dependent kinases (CDKs) regulate a var...
1	1	Abstract Background Non-small cell lung canc...
2	2	Abstract Background Non-small cell lung canc...

Even though we knew its column names which we defined ourselves, we notice (from deeper exploration of datasets) that ID column is common and matches fro both datasets

```
[10]: data_text.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3321 entries, 0 to 3320
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    ID      3321 non-null    int64
1    TEXT     3316 non-null    object
```

```
dtypes: int64(1), object(1)
```

```
memory usage: 52.0+ KB
```

```
data_text has null values
```

```
[11]: # data_text.describe() is not useful for text data values
data_text.columns
```

```
[11]: Index(['ID', 'TEXT'], dtype='object')
```

```
[12]: # checking the dimensions (which we already know)
data_text.shape
```

```
[12]: (3321, 2)
```

```
[13]: # Confirmation of available results
data_variants.Class.unique()
```

```
[13]: array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```

This is discrete data so it is **classification** problem and since there are multiple discrete output possible so we can call it **Multi class** classification problem

Important note : This is medical related problem so correct results are very important. Error can be really costly here so we would like to have result for each class in terms of Probability. We might not be much bothered about time taken by ML algorithm as far as it is reasonable.

We also want our model to be highly interpretable because a medical practitioner wants to also give proper reasoning on why ML algorithm is predicting any class.

We will evaluate our model using Confusion matrix and Multi class log-loss

Data preparation

```
[14]: # We would like to remove all stop words like a, is, an, the, ...
# so we collect all of them from nltk library
stop_words = set(stopwords.words('english'))
```

```
[15]: # defining function to remove all stop words from data
def data_text_preprocess(total_text, ind, col):
    # Remove int values from text data as that might not be imp
    if type(total_text) is not int:
        string = ""
        # replacing all special char with space
        total_text = re.sub('[^a-zA-Z0-9\n]', ' ', str(total_text))
        # replacing multiple spaces with single space
        total_text = re.sub('\s+', ' ', str(total_text))
        # bring whole text to same lower-case scale.
        total_text = total_text.lower()
```

```

for word in total_text.split():
    # if the word is a not a stop word then retain that word from text
    if not word in stop_words:
        string += word + " "

data_text[col][ind] = string

```

```

[16]: # applying data_text_preprocess to data_text
for index, row in data_text.iterrows():
    if type(row['TEXT']) is str:
        data_text_preprocess(row['TEXT'], index, 'TEXT')

```

```

[17]: # merging both gene_variations and text data based on ID
result = pd.merge(data_variants, data_text, on='ID', how='left')
result.head()

```

```

[17]:
   ID  Gene      Variation  Class \
0   0  FAM58A  Truncating Mutations    1
1   1   CBL           W802*         2
2   2   CBL           Q249E         2
3   3   CBL           N454D         3
4   4   CBL           L399V         4

                                     TEXT
0  cyclin dependent kinases cdks regulate variety...
1  abstract background non small cell lung cancer...
2  abstract background non small cell lung cancer...
3  recent evidence demonstrated acquired uniparen...
4  oncogenic mutations monomeric casitas b lineag...

```

```

[18]: # checking for missing values
# missing values may create qualitative problematics in our final analysis
result[result.isnull().any(axis=1)]

```

```

[18]:
   ID  Gene      Variation  Class TEXT
1109 1109  FANCA          S1088F    1  NaN
1277 1277  ARID5B  Truncating Mutations    1  NaN
1407 1407  FGFR3          K508M     6  NaN
1639 1639  FLT1      Amplification     6  NaN
2755 2755  BRAF          G596C     7  NaN

```

We can see many rows with missing data. Now the question is what to do with this missing value. One way could be that we can drop these rows having missing values or we can do some imputation in it. Let's go with imputation only. But question is what to impute here

```

[19]: # Imputing the missing values as "Gene" + "Variation" text

```

```
# This is the example imputation used in the course (check note in the
↳beginning)
result.loc[result['TEXT'].isnull(), 'TEXT'] = result['Gene'] + '
↳'+result['Variation']
```

```
[20]: # Confirming that there are no missing values
result[result.isnull().any(axis=1)]
```

```
[20]: Empty DataFrame
Columns: [ID, Gene, Variation, Class, TEXT]
Index: []
```

Awesome, so all missing values are gone now.

0.2 Creating Training, Test and Validation data

Before we split the data into training, test and validation data set. We want to ensure that all spaces in Gene and Variation column to be replaced by _.

```
[21]: y_true = result['Class'].values
# replacing spaces with "_"
result.Gene = result.Gene.str.replace('\s+', '_')
result.Variation = result.Variation.str.replace('\s+', '_')
```

```
[22]: # Splitting the data into train and test set
X_train, test_df, y_train, y_test = train_test_split(result, y_true,
↳stratify=y_true, test_size=0.2)
# split the train data now into train validation and cross validation
train_df, cv_df, y_train, y_cv = train_test_split(X_train, y_train,
↳stratify=y_train, test_size=0.2)
```

```
[23]: print('Number of data points in train data:', train_df.shape[0])
print('Number of data points in test data:', test_df.shape[0])
print('Number of data points in cross validation data:', cv_df.shape[0])
```

Number of data points in train data: 2124

Number of data points in test data: 665

Number of data points in cross validation data: 532

Checking on train, validation, and test datasets distribution


```
[24]: # counting class values for each set and sorting for better comparison
train_class_distribution = train_df['Class'].value_counts().sort_index(axis=0,
↳level=None, ascending=True, inplace=False, kind='quicksort',
↳na_position='last', sort_remaining=True)
test_class_distribution = test_df['Class'].value_counts().sort_index(axis=0,
↳level=None, ascending=True, inplace=False, kind='quicksort',
↳na_position='last', sort_remaining=True)
```

```
cv_class_distribution = cv_df['Class'].value_counts().sort_index(axis=0,
↳level=None, ascending=True, inplace=False, kind='quicksort',
↳na_position='last', sort_remaining=True)
```

```
[25]: train_class_distribution
```

```
[25]: 1    363
      2    289
      3     57
      4    439
      5    155
      6    176
      7    609
      8     12
      9     24
      Name: Class, dtype: int64
```

```
[26]: # Visualizing train class distrubution
      my_colors = 'rgbkymc'
      train_class_distribution.plot(kind='bar')
      plt.xlabel('Class')
      plt.ylabel(' Number of Data points per Class')
      plt.title('Distribution of yi in train data')
      plt.grid()
      plt.show()
```




**** Visualizing for train class distrubution****

```
[27]: # Printing distribution in percentage form
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ': ', train_class_distribution.
↪values[i], '(', np.round((train_class_distribution.values[i]/train_df.
↪shape[0]*100), 3), '%)')
```

```
Number of data points in class 7 : 609 ( 28.672 %)
Number of data points in class 4 : 439 ( 20.669 %)
Number of data points in class 1 : 363 ( 17.09 %)
Number of data points in class 2 : 289 ( 13.606 %)
Number of data points in class 6 : 176 ( 8.286 %)
Number of data points in class 5 : 155 ( 7.298 %)
Number of data points in class 3 : 57 ( 2.684 %)
Number of data points in class 9 : 24 ( 1.13 %)
Number of data points in class 8 : 12 ( 0.565 %)
```

```
[28]: # test set diribution visualization
my_colors = 'rgbkymc'
test_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Number of Data points per Class')
plt.title('Distribution of yi in test data')
plt.grid()
plt.show()
```





```
[29]: # test set distribution in percentage form
sorted_yi = np.argsort(-test_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ': ', test_class_distribution.
          ↪ values[i], '(', np.round((test_class_distribution.values[i]/test_df.
          ↪ shape[0]*100), 3), '%)')
```

```
Number of data points in class 7 : 191 ( 28.722 %)
Number of data points in class 4 : 137 ( 20.602 %)
Number of data points in class 1 : 114 ( 17.143 %)
Number of data points in class 2 : 91 ( 13.684 %)
Number of data points in class 6 : 55 ( 8.271 %)
Number of data points in class 5 : 48 ( 7.218 %)
Number of data points in class 3 : 18 ( 2.707 %)
Number of data points in class 9 : 7 ( 1.053 %)
Number of data points in class 8 : 4 ( 0.602 %)
```

```
[30]: # cross validation set distribution visualization
my_colors = 'rgbkymc'
cv_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Number of Data points per Class')
plt.title('Distribution of yi in test data')
plt.grid()
```

```
plt.show()
```



```
[31]: # cross validation set distribution in percentage form
sorted_yi = np.argsort(-cv_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ': ', test_class_distribution.
          ↪ values[i], '(', np.round((test_class_distribution.values[i]/test_df.
          ↪ shape[0]*100), 3), '%)')
```

```
Number of data points in class 7 : 191 ( 28.722 %)
Number of data points in class 4 : 137 ( 20.602 %)
Number of data points in class 1 : 114 ( 17.143 %)
Number of data points in class 2 : 91 ( 13.684 %)
Number of data points in class 6 : 55 ( 8.271 %)
Number of data points in class 5 : 48 ( 7.218 %)
Number of data points in class 3 : 18 ( 2.707 %)
Number of data points in class 9 : 7 ( 1.053 %)
Number of data points in class 8 : 4 ( 0.602 %)
```

0.3 Building a Random model - WORST MODEL

```
[32]: # getting the length of our datasets
test_data_len = test_df.shape[0]
cv_data_len = cv_df.shape[0]

# creating an output array that has exactly same size as the CV data
cv_predicted_y = np.zeros((cv_data_len,9))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,9)
    # setting random values so each row sum is equal to 1
    cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
# Evaluating and printing the Log Loss for worst model
# All models that will be generated should not be worse than worst model
print("Log loss on Cross Validation Data using Random_
↪Model",log_loss(y_cv,cv_predicted_y, eps=1e-15))

# Test-Set error (worst model).
# creating output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,9))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,9)
    test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random_
↪Model",log_loss(y_test,test_predicted_y, eps=1e-15))
```

Log loss on Cross Validation Data using Random Model 2.4983107431467615
Log loss on Test Data using Random Model 2.4450267324737074

```
[33]: # Lets get the index of max probablity
predicted_y =np.argmax(test_predicted_y, axis=1)
predicted_y
```

```
[33]: array([4, 7, 3, 5, 6, 2, 1, 1, 1, 5, 6, 8, 5, 1, 8, 0, 3, 5, 6, 6, 0, 3,
        8, 8, 6, 5, 7, 1, 5, 4, 0, 8, 7, 7, 5, 4, 7, 5, 4, 1, 1, 8, 8, 8,
        6, 8, 7, 1, 6, 7, 0, 0, 2, 5, 4, 2, 1, 0, 4, 8, 1, 2, 7, 5, 5, 1,
        2, 3, 0, 6, 7, 8, 5, 4, 8, 2, 1, 6, 6, 7, 2, 2, 8, 0, 7, 2, 7, 1,
        4, 8, 1, 2, 3, 6, 6, 6, 4, 2, 1, 7, 1, 2, 2, 2, 4, 8, 1, 5, 5, 3,
        2, 1, 0, 0, 1, 7, 4, 4, 2, 1, 0, 1, 6, 6, 7, 7, 2, 0, 1, 7, 5, 3,
        6, 0, 7, 8, 4, 5, 5, 6, 6, 2, 6, 0, 1, 5, 4, 0, 4, 6, 7, 1, 6, 6,
        8, 7, 4, 8, 6, 4, 8, 2, 7, 8, 1, 1, 6, 6, 2, 2, 8, 6, 8, 4, 2, 8,
        0, 3, 6, 8, 0, 7, 5, 0, 1, 7, 7, 7, 3, 1, 2, 1, 8, 8, 2, 0, 3, 3,
        7, 8, 3, 0, 6, 3, 6, 7, 0, 4, 3, 8, 2, 6, 8, 7, 7, 4, 8, 5, 8, 1,
        3, 6, 0, 5, 5, 7, 7, 0, 6, 3, 5, 0, 5, 6, 2, 7, 1, 5, 4, 8, 8, 5,
        1, 7, 4, 7, 6, 0, 2, 1, 5, 2, 7, 3, 6, 3, 5, 2, 2, 5, 8, 4, 4, 6,
        1, 5, 2, 2, 5, 0, 3, 4, 8, 6, 3, 2, 8, 4, 8, 0, 5, 0, 2, 5, 8, 4,
        7, 5, 8, 5, 7, 8, 5, 7, 8, 0, 5, 4, 7, 4, 0, 0, 6, 4, 0, 7, 5, 3,
```

```

7, 1, 4, 7, 3, 4, 3, 7, 8, 3, 3, 8, 0, 3, 2, 3, 5, 1, 1, 6, 5, 7,
7, 0, 5, 0, 8, 0, 1, 3, 8, 8, 7, 6, 6, 6, 3, 2, 4, 4, 5, 6, 4, 2,
3, 3, 2, 0, 1, 2, 3, 3, 6, 5, 6, 1, 3, 7, 2, 4, 0, 7, 2, 7, 1, 3,
8, 2, 2, 2, 3, 4, 1, 1, 3, 1, 2, 6, 1, 6, 0, 6, 7, 2, 4, 0, 5, 5,
3, 4, 6, 6, 4, 4, 3, 2, 8, 5, 4, 4, 2, 0, 3, 6, 7, 1, 1, 0, 7, 0,
8, 5, 2, 0, 5, 1, 7, 4, 8, 0, 5, 1, 5, 2, 8, 1, 0, 3, 1, 5, 4, 6,
2, 5, 5, 4, 1, 2, 2, 2, 0, 6, 6, 2, 6, 2, 7, 6, 5, 2, 2, 2, 2, 2,
5, 7, 0, 6, 7, 6, 6, 2, 4, 0, 8, 0, 2, 6, 7, 5, 5, 0, 5, 7, 5, 1,
1, 8, 4, 5, 8, 6, 3, 2, 7, 8, 5, 6, 6, 3, 4, 2, 5, 5, 0, 1, 5, 3,
6, 4, 2, 0, 6, 7, 1, 4, 6, 7, 5, 8, 4, 3, 0, 8, 0, 2, 6, 2, 2, 2,
1, 1, 1, 0, 4, 6, 4, 4, 2, 0, 4, 8, 8, 4, 7, 5, 7, 3, 2, 3, 4, 8,
2, 6, 6, 8, 4, 3, 8, 7, 1, 1, 3, 1, 8, 1, 2, 8, 3, 6, 8, 0, 7, 0,
5, 7, 7, 6, 2, 7, 6, 3, 6, 5, 6, 2, 7, 0, 0, 1, 1, 5, 4, 1, 0, 3,
0, 2, 8, 6, 3, 3, 6, 6, 6, 3, 3, 5, 6, 0, 2, 7, 3, 5, 7, 8, 7, 4,
3, 7, 6, 7, 1, 5, 3, 4, 2, 5, 0, 2, 5, 2, 6, 2, 5, 7, 3, 8, 2, 7,
3, 5, 1, 0, 3, 6, 0, 8, 3, 7, 4, 8, 0, 8, 3, 4, 7, 6, 6, 5, 8, 1,
4, 6, 5, 6, 5])

```

The index value ranging from 0 to 8. Each value should be increased by 1.


```
[34]: predicted_y = predicted_y + 1
```

Confusion Matrix

```
[35]: C = confusion_matrix(y_test, predicted_y)
C
```

```
[35]: array([[14,  9, 15,  8, 20,  8, 13, 12, 15],
            [10,  8, 12, 16,  6, 12, 10,  8,  9],
            [ 3,  3,  1,  2,  1,  3,  1,  1,  3],
            [10, 15, 18, 13, 10, 19, 18, 17, 17],
            [ 7,  4,  4,  4,  6,  3,  8,  7,  5],
            [ 8,  8,  8,  4,  3,  8,  9,  6,  1],
            [13, 21, 25, 16, 18, 26, 27, 24, 21],
            [ 2,  1,  0,  0,  0,  0,  0,  1,  0],
            [ 1,  0,  0,  1,  1,  0,  1,  2,  1]])
```


```
[36]: # Displaying predictions
labels = [1,2,3,4,5,6,7,8,9]
plt.figure(figsize=(20,7))
sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
            yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



The diagonal of confusion matrix gives the correctly predicted values (values with class 1 predicted as class 1 for the first cell of example)

Precision matrix

```
[37]: # generating precision matrix
B =(C/C.sum(axis=0))
# displaying precision matrix
plt.figure(figsize=(20,7))
sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```




Precision matrix is same as confusion matrix put cells are expressed as percentages in the same column (values predicted as class Xi) Cell 1,1 shows how many percent of points predicted as Class

1 are actually Class 1

Recall matrix

```
[38]: # generating recall matrix
A = (((C.T)/(C.sum(axis=1))).T)
# displaying recall matrix
plt.figure(figsize=(20,7))
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
            yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()
```



recall matrix is same as confusion matrix but values are expressed as percentage in row direction meaning each cell represents how percent of class *i* are predicted as class *j*

0.4 Evaluating Gene Column

Now we will look at each independent column to make sure its relevant for my target variable but the question is, how? Let's understand with our first column Gene which is categorical in nature.

So, let's explore column **Gene** and let's look at its distribution.

```
[39]: # Creating unique genes series
unique_genes = train_df['Gene'].value_counts()
# measuring the length of unique genes matrix
print('Number of Unique Genes :', unique_genes.shape[0])
# the top 10 genes that occurred most
print(unique_genes.head(10))
```


Number of Unique Genes : 225

BRCA1 171

TP53 102

```
EGFR      98
PTEN      81
BRCA2     78
KIT       60
BRAF      57
ERBB2     45
ALK       43
PDGFRA    41
Name: Gene, dtype: int64
```

```
[40]: # Cumulative distribution of unique genes
s = sum(unique_genes.values);
h = unique_genes.values/s;
c = np.cumsum(h)
plt.plot(c,label='Cumulative distribution of Genes')
plt.grid()
plt.legend()
plt.show()
```



From the graph - Top 50 Gene attributes make almost 75% of all data

We need to convert these categorical variable to appropriate format which my machine learning algorithm will be able to take as an input.

So we have 2 techniques to deal with it.

One-hot encoding

Response Encoding (Mean imputation)

We will use both of them to see which one work the best. So lets start encoding using one hot encoder For Response encoding we will add columns representing class values filled with values of probability of that class given GENE - i

One hot encoder of genes classification

```
[41]: # one-hot encoding of Gene feature.
gene_vectorizer = CountVectorizer()
train_gene_feature_onehotCoding = gene_vectorizer.
    ↪fit_transform(train_df['Gene'])
test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Gene'])
cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene'])
```

```
[42]: # checking classified train set shape
train_gene_feature_onehotCoding.shape
```

```
[42]: (2124, 225)
```

```
[43]: # plotting added column names
gene_vectorizer.get_feature_names()
```

```
[43]: ['abl1',
      'acvr1',
      'ago2',
      'akt1',
      'akt2',
      'akt3',
      'alk',
      'apc',
      'ar',
      'araf',
      'arid1a',
      'arid1b',
      'arid5b',
      'asxl2',
      'atm',
      'atrx',
      'aurka',
      'aurkb',
      'axin1',
      'b2m',
      'bap1',
      'bcl2l11',
      'bcor',
      'braf',
      'brca1',
      'brca2',
      'brd4',
```


'brip1',
'btk',
'card11',
'carm1',
'casp8',
'cbl',
'ccnd1',
'ccnd3',
'ccne1',
'cdh1',
'cdk12',
'cdk4',
'cdk6',
'cdk8',
'cdkn1a',
'cdkn1b',
'cdkn2a',
'cdkn2b',
'chek2',
'cic',
'crebbp',
'ctcf',
'ctla4',
'ctnnb1',
'ddr2',
'dicer1',
'dnmt3a',
'dnmt3b',
'egfr',
'eif1ax',
'elf3',
'ep300',
'epas1',
'epcam',
'erbb2',
'erbb3',
'erbb4',
'ercc2',
'ercc4',
'erg',
'errfi1',
'esr1',
'etv1',
'etv6',
'ewsr1',
'ezh2',
'fanca',

'fance',
'fat1',
'fbxw7',
'fgf19',
'fgfr1',
'fgfr2',
'fgfr3',
'fgfr4',
'flt1',
'flt3',
'foxa1',
'foxl2',
'foxo1',
'fubp1',
'gata3',
'gli1',
'gnaq',
'gnas',
'h3f3a',
'hist1h1c',
'hla',
'hnf1a',
'hras',
'idh1',
'idh2',
'igf1r',
'ikzf1',
'il7r',
'inpp4b',
'jak1',
'jak2',
'jun',
'kdm5a',
'kdm5c',
'kdm6a',
'kdr',
'keap1',
'kit',
'klf4',
'kmt2a',
'kmt2c',
'kmt2d',
'knstrn',
'kras',
'map2k1',
'map2k2',
'map2k4',

'map3k1',
'mapk1',
'mdm4',
'med12',
'mef2b',
'men1',
'met',
'mga',
'mlh1',
'mpl',
'msh2',
'msh6',
'mtor',
'myc',
'mycn',
'myd88',
'nf1',
'nf2',
'nfe2l2',
'nfkb1a',
'nkx2',
'notch1',
'notch2',
'npm1',
'nras',
'nsd1',
'ntrk1',
'ntrk2',
'ntrk3',
'nup93',
'pak1',
'pbrm1',
'pdgfra',
'pdgfrb',
'pik3ca',
'pik3cb',
'pik3cd',
'pik3r1',
'pik3r2',
'pim1',
'pms1',
'pms2',
'pole',
'ppp2r1a',
'ppp6c',
'prdm1',
'ptch1',

'pten',
'ptpn11',
'ptprd',
'ptprt',
'rab35',
'rac1',
'rad21',
'rad50',
'rad51b',
'raf1',
'rasa1',
'rb1',
'rbm10',
'ret',
'rheb',
'rhoa',
'rit1',
'ros1',
'runx1',
'rxra',
'rybp',
'sdhb',
'sdhc',
'setd2',
'sf3b1',
'shoc2',
'shq1',
'smad2',
'smad3',
'smad4',
'smarca4',
'smarcb1',
'smo',
'sos1',
'sox9',
'spop',
'src',
'srsf2',
'stag2',
'stat3',
'stk11',
'tcf3',
'tert',
'tet1',
'tet2',
'tgfbr1',
'tmprss2',

```

'tp53',
'tsc1',
'tsc2',
'u2af1',
'vegfa',
'vhl',
'whsc1',
'xpo1',
'xrcc2',
'yap1']

```

**** Response Encoding of genes classification**** Code is as given in the course (i do not fully understand Response Encoding yet). Course link given in description at the begining of this notebook

```

[44]: # ----Notes----
# code for response coding with Laplace smoothing.
# alpha : used for laplace smoothing
# feature: ['gene', 'variation']
# df: ['train_df', 'test_df', 'cv_df']
# algorithm
# -----
# Consider all unique values and the number of occurances of given feature in
→train data dataframe
# build a vector (1*9) , the first element = (number of times it occurred in
→class1 + 10*alpha / number of time it occurred in total data+90*alpha)
# gv_dict is like a look up table, for every gene it store a (1*9)
→representation of it
# for a value of feature in df:
# if it is in train data:
# we add the vector that was stored in 'gv_dict' look up table to 'gv_fea'
# if it is not there is train:
# we add [1/9, 1/9, 1/9, 1/9,1/9, 1/9, 1/9, 1/9, 1/9] to 'gv_fea'
# return 'gv_fea'
# -----

# get_gv_fea_dict: Get Gene varaition Feature Dict
def get_gv_fea_dict(alpha, feature, df):
    # value_count: it contains a dict like
    # print(train_df['Gene'].value_counts())
    # output:
    #          {BRCA1      174
    #           TP53      106
    #           EGFR       86
    #           BRCA2       75
    #           PTEN       69

```

```

#         ...}
# print(train_df['Variation'].value_counts())
# output:
# {
# Truncating_Mutations          63
# Deletion                     43
# Amplification                 43
# Fusions                      22
# Overexpression                3
# E17K                         3
# Q61L                         3
# S222D                        2
# P130S                        2
# ...
# }
value_count = train_df[feature].value_counts()

# gv_dict : Gene Variation Dict, which contains the probability array for
→ each gene/variation
gv_dict = dict()

# denominator will contain the number of time that particular feature
→ occurred in whole data
for i, denominator in value_count.items():
    # vec will contain (p(yi==1/Gi) probability of gene/variation belongs
→ to particular class
    # vec is 9 dimensional vector
    vec = []
    for k in range(1,10):
        # print(train_df.loc[(train_df['Class']==1) &
→ (train_df['Gene']=='BRCA1')])
        #
        # ID   Gene      Variation  Class
        # 2470 2470  BRCA1      S1715C    1
        # 2486 2486  BRCA1      S1841R    1
        # 2614 2614  BRCA1        M1R     1
        # 2432 2432  BRCA1      L1657P    1
        # 2567 2567  BRCA1      T1685A    1
        # 2583 2583  BRCA1      E1660G    1
        # 2634 2634  BRCA1      W1718L    1
        # cls_cnt.shape[0] will return the number of rows

        cls_cnt = train_df.loc[(train_df['Class']==k) &
→ (train_df[feature]==i)]

        # cls_cnt.shape[0](numerator) will contain the number of time that
→ particular feature occurred in whole data

```

```

        vec.append((cls_cnt.shape[0] + alpha*10)/ (denominator + 90*alpha))

        # we are adding the gene/variation to the dict as key and vec as value
        gv_dict[i]=vec
    return gv_dict

# Get Gene variation feature
def get_gv_feature(alpha, feature, df):
    # print(gv_dict)
    # {'BRCA1': [0.20075757575757575, 0.03787878787878788, 0.
    → 068181818181818177, 0.13636363636363635, 0.25, 0.19318181818181818, 0.
    → 03787878787878788, 0.03787878787878788, 0.03787878787878788],
    # 'TP53': [0.32142857142857145, 0.061224489795918366, 0.
    → 061224489795918366, 0.27040816326530615, 0.061224489795918366, 0.
    → 066326530612244902, 0.051020408163265307, 0.051020408163265307, 0.
    → 056122448979591837],
    # 'EGFR': [0.056818181818181816, 0.21590909090909091, 0.0625, 0.
    → 068181818181818177, 0.068181818181818177, 0.0625, 0.34659090909090912, 0.
    → 0625, 0.056818181818181816],
    # 'BRCA2': [0.13333333333333333, 0.060606060606060608, 0.
    → 060606060606060608, 0.078787878787878782, 0.13939393939393934, 0.
    → 34545454545454546, 0.060606060606060608, 0.060606060606060608, 0.
    → 060606060606060608],
    # 'PTEN': [0.069182389937106917, 0.062893081761006289, 0.
    → 069182389937106917, 0.46540880503144655, 0.075471698113207544, 0.
    → 062893081761006289, 0.069182389937106917, 0.062893081761006289, 0.
    → 062893081761006289],
    # 'KIT': [0.066225165562913912, 0.25165562913907286, 0.
    → 072847682119205295, 0.072847682119205295, 0.066225165562913912, 0.
    → 066225165562913912, 0.27152317880794702, 0.066225165562913912, 0.
    → 066225165562913912],
    # 'BRAF': [0.066666666666666666, 0.17999999999999999, 0.
    → 073333333333333334, 0.073333333333333334, 0.09333333333333338, 0.
    → 080000000000000002, 0.29999999999999999, 0.066666666666666666, 0.
    → 066666666666666666],
    # ...
    # }
    gv_dict = get_gv_fea_dict(alpha, feature, df)
    # value_count is similar in get_gv_fea_dict
    value_count = train_df[feature].value_counts()

    # gv_fea: Gene_variation feature, it will contain the feature for each
    → feature value in the data
    gv_fea = []
    # for every feature values in the given data frame we will check if it is
    → there in the train data then we will add the feature to gv_fea

```

```

# if not we will add [1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9] to gv_fea
for index, row in df.iterrows():
    if row[feature] in dict(value_count).keys():
        gv_fea.append(gv_dict[row[feature]])
    else:
        gv_fea.append([1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9])
#
    gv_fea.append([-1,-1,-1,-1,-1,-1,-1,-1,-1])
return gv_fea

```

```

[45]: # response-coding of the Gene feature
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene",
    ↪train_df))
# test gene feature
test_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene",
    ↪test_df))
# cross validation gene feature
cv_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", cv_df))

```

```

[46]: # exploring training gene shape
train_gene_feature_responseCoding.shape

```

```

[46]: (2124, 9)

```

Response encoding adds a number of classes equal to number of classes (possible results)

We will build model having only gene column with one hot encoder with simple model like Logistic regression. If log loss with only one column Gene comes out to be better than random model, than this feature is important.

```

[47]: # We need a hyperparameter for SGD classifier.
# giving alpha a set of ranges to compare
alpha = [10 ** x for x in range(-5, 1)]

```

```

[48]: # We will be using SGD classifier
# http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.
    ↪SGDClassifier.html
# We will also be using Calibrated Classifier to get the result into probability
    ↪format t be used for log loss
cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_gene_feature_onehotCoding, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_gene_feature_onehotCoding, y_train)

```



```

    predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_,
    ↪eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv,
    ↪predict_y, labels=clf.classes_, eps=1e-15))

```

```

For values of alpha = 1e-05 The log loss is: 1.1988978880599888
For values of alpha = 0.0001 The log loss is: 1.1907921428324062
For values of alpha = 0.001 The log loss is: 1.251415171375968
For values of alpha = 0.01 The log loss is: 1.3661590369405565
For values of alpha = 0.1 The log loss is: 1.4194498089685597
For values of alpha = 1 The log loss is: 1.448480654741542

```

```

[49]: # Lets plot the same to check the best Alpha value
      #fig, ax = plt.subplots()
      #ax.plot(alpha, cv_log_error_array, c='g')
      #for i, txt in enumerate(np.round(cv_log_error_array,3)):
      #    ax.annotate((alpha[i], np.round(txt,3)), (alpha[i], cv_log_error_array[i]))
      #plt.grid()
      #plt.title("Cross Validation Error for each alpha")
      #plt.xlabel("Alpha i's")
      #plt.ylabel("Error measure")
      #plt.show()

      # because the graph is not clear this section is commented and written to just
      ↪give the best alpha
      def print_best_alpha(alpha_arr, loss_arr):
          print("The best alpha is: " + str(alpha_arr[loss_arr.index(min(loss_arr))])
          ↪+ "\n\nThe best alpha index is: " + str(loss_arr.index(min(loss_arr))))
      print_best_alpha(alpha, cv_log_error_array)

```

```

The best alpha is: 0.0001
The best alpha index is: 1

```

```

[50]: # Lets use best alpha value as we can see from above graph and compute log loss
      # Building a very simple model using just gene column to check error decrease
      ↪from worst model
      # building a model with just one feature gives information on how much that
      ↪feature is meaningful for the final result
      best_alpha = np.argmin(cv_log_error_array)
      clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
      ↪random_state=42)
      clf.fit(train_gene_feature_onehotCoding, y_train)
      sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
      sig_clf.fit(train_gene_feature_onehotCoding, y_train)

```

```

predict_y = sig_clf.predict_proba(train_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
↪",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation_
↪log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
↪",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

For values of best alpha = 0.0001 The train log loss is: 0.9795845316837705
 For values of best alpha = 0.0001 The cross validation log loss is:
 1.1907921428324062
 For values of best alpha = 0.0001 The test log loss is: 1.238282196746077

Now lets check how many values are overlapping between train, test or between CV and train

```

[51]: # checking for overlapping between train set and [cross_validation, test] set
test_coverage=test_df[test_df['Gene'].isin(list(set(train_df['Gene'])))].
↪shape[0]
cv_coverage=cv_df[cv_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]

[52]: print('1. In test data',test_coverage, 'out of',test_df.shape[0], ":
↪", (test_coverage/test_df.shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0],":")
↪, (cv_coverage/cv_df.shape[0])*100)

```

1. In test data 632 out of 665 : 95.03759398496241
 2. In cross validation data 515 out of 532 : 96.80451127819549

Overlapping of training set with cross-validation and test set is very high which is very good for our models

0.5 Evaluating Variation column

Variation is also a categorical variable so we have to deal in same way like we have done for *Gene* column. We will again get the one hot encoder and response encoding variable for variation column.

```

[53]: unique_variations = train_df['Variation'].value_counts()
print('Number of Unique Variations :', unique_variations.shape[0])
# the top 10 variations that occurred most
print(unique_variations.head(10))


```

Number of Unique Variations : 1952
 Truncating_Mutations 57
 Deletion 46
 Amplification 36
 Fusions 16
 Overexpression 5

```
T58I          3
G13C          2
ETV6-NTRK3_Fusion  2
Q61H          2
F28L          2
Name: Variation, dtype: int64
```

```
[54]: # looking at the cumulative distribution of unique variation values
s = sum(unique_variations.values);
h = unique_variations.values/s;
c = np.cumsum(h)
print(c)
plt.plot(c,label='Cumulative distribution of Variations')
plt.grid()
plt.legend()
plt.show()
```

```
[0.02683616 0.04849341 0.06544256 ... 0.99905838 0.99952919 1.          ]
```



0.6 Evaluating Variation column

Variation is also a categorical variable so we have to deal in same way like we have done for Gene column. We will again get the one hot encoder and response encoding variable for variation column.

```
[55]: unique_variations = train_df['Variation'].value_counts()
print('Number of Unique Variations :', unique_variations.shape[0])
```

```
# the top 10 variations that occurred most
print(unique_variations.head(10))
```


Number of Unique Variations : 1952

Truncating_Mutations	57
Deletion	46
Amplification	36
Fusions	16
Overexpression	5
T58I	3
G13C	2
ETV6-NTRK3_Fusion	2
Q61H	2
F28L	2

Name: Variation, dtype: int64

```
[56]: # plotting the distribution of variation values
s = sum(unique_variations.values);
h = unique_variations.values/s;
c = np.cumsum(h)
print(c)
plt.plot(c,label='Cumulative distribution of Variations')
plt.grid()
plt.legend()
plt.show()
```

```
[0.02683616 0.04849341 0.06544256 ... 0.99905838 0.99952919 1.      ]
```



```
[57]: # one-hot encoding of variation values.
variation_vectorizer = CountVectorizer()
train_variation_feature_onehotCoding = variation_vectorizer.
    ↪fit_transform(train_df['Variation'])
test_variation_feature_onehotCoding = variation_vectorizer.
    ↪transform(test_df['Variation'])
cv_variation_feature_onehotCoding = variation_vectorizer.
    ↪transform(cv_df['Variation'])
```

```
[58]: # The shape of one hot encoder column for variation
train_variation_feature_onehotCoding.shape
```

```
[58]: (2124, 1976)
```

```
[59]: # Response encoding of variation values.
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_variation_feature_responseCoding = np.array(get_gv_feature(alpha, ↵
    ↪"Variation", train_df))
# test gene feature
test_variation_feature_responseCoding = np.array(get_gv_feature(alpha, ↵
    ↪"Variation", test_df))
# cross validation gene feature
cv_variation_feature_responseCoding = np.array(get_gv_feature(alpha, ↵
    ↪"Variation", cv_df))
```

```
[60]: # the shape of this response encoding result
train_variation_feature_responseCoding.shape
```

```
[60]: (2124, 9)
```

```
[61]: # Lets again build the model with only column name of variation column
# We need a hyperparameter for SGD classifier.
alpha = [10 ** x for x in range(-5, 1)]
# We will be using SGD classifier
# http://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.
    ↪SGDClassifier.html
# We will also be using Calibrated Classifier to get the result into probability ↵
    ↪format t be used for log loss
cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_variation_feature_onehotCoding, y_train)
```

```

sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_variation_feature_onehotCoding, y_train)
predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)

cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_,
↪eps=1e-15))
print('For values of alpha = ', i, "The log loss is:",log_loss(y_cv,
↪predict_y, labels=clf.classes_, eps=1e-15))

```

```

For values of alpha = 1e-05 The log loss is: 1.7141045692201116
For values of alpha = 0.0001 The log loss is: 1.709502946278758
For values of alpha = 0.001 The log loss is: 1.7173042028158032
For values of alpha = 0.01 The log loss is: 1.7275108511038793
For values of alpha = 0.1 The log loss is: 1.7363861296986798
For values of alpha = 1 The log loss is: 1.73757998584688

```

```
[62]: print_best_alpha(alpha, cv_log_error_array)
```

```

The best alpha is: 0.0001
The best alpha index is: 1

```

```

[63]: # checking for error on a simple model buided using just variation column
best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
↪random_state=42)
clf.fit(train_variation_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_variation_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
↪",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
↪log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
↪",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

For values of best alpha = 0.0001 The train log loss is: 0.6085583338786402
For values of best alpha = 0.0001 The cross validation log loss is:
1.709502946278758
For values of best alpha = 0.0001 The test log loss is: 1.6929151540351735

```

There is very high difference on error on the training set with errors in cross-validation and test set

```
[64]: # checking overlapping of training set with other sets
```

```

test_coverage=test_df[test_df['Variation']].
    ↳isin(list(set(train_df['Variation'])))].shape[0]
cv_coverage=cv_df[cv_df['Variation'].isin(list(set(train_df['Variation'])))].
    ↳shape[0]
print('1. In test data',test_coverage, 'out of',test_df.shape[0], ":
    ↳", (test_coverage/test_df.shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0],": "
    ↳, (cv_coverage/cv_df.shape[0])*100)

```

```

1. In test data 93 out of 665 : 13.984962406015036
2. In cross validation data 53 out of 532 : 9.962406015037594

```

There is very low overlapping. The low overlapping is the main reason for error increase from training set to test and cross-validation set This feature may be considered to be dropped

0.7 Evaluating text column

text column is little different from other columns because it contains many words in a single cell we have already removed all stop words from our text data cells

```

[65]: # cls_text is a data frame
# for every row in data frame consider the 'TEXT'
# split the words by space
# make a dict with those words
# increment its count whenever we see that word

def extract_dictionary_paddle(cls_text):
    dictionary = defaultdict(int)
    for index, row in cls_text.iterrows():
        for word in row['TEXT'].split():
            dictionary[word] +=1
    return dictionary

#https://stackoverflow.com/a/1602964
def get_text_responsecoding(df):
    text_feature_responseCoding = np.zeros((df.shape[0],9))
    for i in range(0,9):
        row_index = 0
        for index, row in df.iterrows():
            sum_prob = 0
            for word in row['TEXT'].split():
                sum_prob += math.log(((dict_list[i].get(word,0)+10 )/
    ↳(total_dict.get(word,0)+90)))
                text_feature_responseCoding[row_index][i] = math.exp(sum_prob/
    ↳len(row['TEXT'].split()))
            row_index += 1
    return text_feature_responseCoding

```

```
[66]: # building a CountVectorizer with all the words that occurred minimum 3 times in
      ↪ train data
text_vectorizer = CountVectorizer(min_df=3)
train_text_feature_onehotCoding = text_vectorizer.
      ↪ fit_transform(train_df['TEXT'])
# getting all the feature names (words)
train_text_features= text_vectorizer.get_feature_names()

# train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and returns
      ↪ (1*number of features) vector
train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1

# zip(list(text_features),text_fea_counts) will zip a word with its number of
      ↪ times it occurred
text_fea_dict = dict(zip(list(train_text_features),train_text_fea_counts))

print("Total number of unique words in train data :", len(train_text_features))
```

Total number of unique words in train data : 53256

```
[67]: dict_list = []
# dict_list=[] contains 9 dictionaries each corresponds to a class
for i in range(1,10):
    cls_text = train_df[train_df['Class']==i]
    # build a word dict based on the words in that class
    dict_list.append(extract_dictionary_paddle(cls_text))
    # append it to dict_list

# dict_list[i] is build on i'th class text data
# total_dict is build on whole training text data
total_dict = extract_dictionary_paddle(train_df)

confuse_array = []
for i in train_text_features:
    ratios = []
    max_val = -1
    for j in range(0,9):
        ratios.append((dict_list[j][i]+10)/(total_dict[i]+90))
    confuse_array.append(ratios)
confuse_array = np.array(confuse_array)
```

```
[68]: # response coding of text features
# text column calculations take a long time
train_text_feature_responseCoding = get_text_responsecoding(train_df)
test_text_feature_responseCoding = get_text_responsecoding(test_df)
cv_text_feature_responseCoding = get_text_responsecoding(cv_df)
```



```
[69]: # https://stackoverflow.com/a/16202486
# we convert each row values such that they sum to 1
train_text_feature_responseCoding = (train_text_feature_responseCoding.T/
    ↪train_text_feature_responseCoding.sum(axis=1)).T
test_text_feature_responseCoding = (test_text_feature_responseCoding.T/
    ↪test_text_feature_responseCoding.sum(axis=1)).T
cv_text_feature_responseCoding = (cv_text_feature_responseCoding.T/
    ↪cv_text_feature_responseCoding.sum(axis=1)).T

[70]: # don't forget to normalize every feature
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding,
    ↪axis=0)

# we use the same vectorizer that was trained on train data
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
# don't forget to normalize every feature
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding,
    ↪axis=0)

# we use the same vectorizer that was trained on train data
cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)

[71]: #https://stackoverflow.com/a/2258273/4084039
sorted_text_fea_dict = dict(sorted(text_fea_dict.items(), key=lambda x: x[1] ,
    ↪reverse=True))
sorted_text_occur = np.array(list(sorted_text_fea_dict.values()))

[72]: # Number of words for a given frequency.
print(Counter(sorted_text_occur))
```

```
Counter({3: 5613, 4: 3647, 6: 2916, 5: 2765, 8: 2326, 7: 2053, 12: 1642, 9:
1593, 10: 1355, 11: 980, 14: 929, 13: 874, 15: 868, 16: 793, 18: 713, 24: 605,
17: 583, 20: 569, 21: 479, 19: 463, 28: 405, 27: 398, 42: 391, 30: 376, 22: 376,
25: 358, 26: 348, 23: 341, 32: 308, 36: 283, 29: 282, 53: 267, 31: 267, 35: 255,
34: 254, 33: 246, 38: 221, 40: 215, 39: 212, 48: 194, 37: 192, 44: 175, 50: 174,
43: 170, 45: 166, 54: 160, 41: 158, 51: 153, 47: 152, 46: 152, 63: 149, 49: 144,
56: 143, 60: 136, 57: 134, 58: 129, 52: 123, 76: 119, 72: 119, 55: 116, 62: 115,
64: 106, 65: 105, 61: 104, 59: 103, 68: 98, 70: 97, 66: 97, 84: 95, 80: 87, 67:
86, 74: 85, 69: 85, 90: 81, 77: 81, 79: 80, 83: 77, 73: 77, 96: 74, 78: 74, 75:
72, 82: 71, 93: 68, 71: 67, 91: 65, 86: 65, 105: 64, 99: 64, 81: 62, 106: 60,
92: 60, 85: 59, 112: 58, 95: 58, 114: 57, 94: 57, 87: 57, 100: 56, 104: 55, 101:
55, 98: 54, 88: 53, 107: 51, 115: 49, 111: 49, 97: 49, 124: 48, 120: 48, 103:
48, 134: 47, 108: 47, 128: 46, 117: 46, 144: 45, 109: 44, 102: 44, 89: 44, 113:
43, 130: 42, 110: 42, 119: 41, 126: 39, 125: 39, 143: 38, 141: 38, 122: 38, 147:
36, 139: 36, 127: 36, 140: 35, 137: 34, 123: 34, 118: 34, 185: 33, 135: 33, 133:
```

33, 121: 33, 116: 33, 240: 32, 159: 32, 148: 32, 142: 32, 131: 32, 150: 31, 145:
 31, 136: 31, 210: 30, 160: 30, 158: 30, 152: 30, 138: 29, 173: 28, 170: 28, 168:
 28, 155: 28, 153: 28, 151: 28, 146: 28, 180: 27, 156: 27, 179: 26, 166: 26, 132:
 26, 212: 25, 183: 25, 175: 25, 167: 25, 154: 25, 202: 24, 191: 24, 190: 24, 162:
 24, 129: 24, 244: 23, 200: 23, 192: 23, 189: 23, 182: 23, 157: 23, 208: 22, 171:
 22, 272: 21, 227: 21, 222: 21, 215: 21, 204: 21, 169: 21, 161: 21, 149: 21, 260:
 20, 256: 20, 216: 20, 195: 20, 186: 20, 181: 20, 178: 20, 165: 20, 163: 20, 242:
 19, 233: 19, 219: 19, 206: 19, 203: 19, 193: 19, 174: 19, 172: 19, 269: 18, 265:
 18, 261: 18, 232: 18, 226: 18, 211: 18, 207: 18, 310: 17, 291: 17, 290: 17, 266:
 17, 245: 17, 239: 17, 229: 17, 217: 17, 214: 17, 196: 17, 187: 17, 176: 17, 284:
 16, 280: 16, 273: 16, 250: 16, 237: 16, 236: 16, 225: 16, 224: 16, 221: 16, 218:
 16, 205: 16, 201: 16, 198: 16, 188: 16, 299: 15, 298: 15, 289: 15, 288: 15, 274:
 15, 246: 15, 241: 15, 230: 15, 228: 15, 184: 15, 164: 15, 400: 14, 397: 14, 292:
 14, 264: 14, 249: 14, 231: 14, 220: 14, 213: 14, 199: 14, 177: 14, 372: 13, 354:
 13, 309: 13, 308: 13, 283: 13, 282: 13, 281: 13, 263: 13, 253: 13, 252: 13, 247:
 13, 243: 13, 235: 13, 384: 12, 370: 12, 340: 12, 327: 12, 319: 12, 302: 12, 276:
 12, 275: 12, 270: 12, 257: 12, 255: 12, 251: 12, 248: 12, 238: 12, 234: 12, 223:
 12, 197: 12, 432: 11, 410: 11, 396: 11, 376: 11, 374: 11, 356: 11, 350: 11, 345:
 11, 342: 11, 338: 11, 333: 11, 324: 11, 320: 11, 317: 11, 314: 11, 301: 11, 287:
 11, 277: 11, 268: 11, 254: 11, 194: 11, 407: 10, 402: 10, 393: 10, 369: 10, 351:
 10, 346: 10, 336: 10, 326: 10, 318: 10, 315: 10, 313: 10, 304: 10, 303: 10, 296:
 10, 293: 10, 285: 10, 279: 10, 262: 10, 259: 10, 209: 10, 756: 9, 646: 9, 557:
 9, 549: 9, 488: 9, 481: 9, 439: 9, 435: 9, 431: 9, 412: 9, 380: 9, 377: 9, 373:
 9, 368: 9, 362: 9, 352: 9, 337: 9, 328: 9, 322: 9, 316: 9, 312: 9, 300: 9, 278:
 9, 271: 9, 648: 8, 573: 8, 541: 8, 539: 8, 535: 8, 516: 8, 494: 8, 461: 8, 453:
 8, 452: 8, 442: 8, 430: 8, 429: 8, 424: 8, 421: 8, 417: 8, 405: 8, 385: 8, 375:
 8, 371: 8, 364: 8, 360: 8, 357: 8, 349: 8, 347: 8, 334: 8, 325: 8, 286: 8, 267:
 8, 258: 8, 754: 7, 591: 7, 567: 7, 536: 7, 508: 7, 507: 7, 484: 7, 471: 7, 467:
 7, 438: 7, 418: 7, 416: 7, 408: 7, 406: 7, 395: 7, 394: 7, 392: 7, 381: 7, 363:
 7, 361: 7, 355: 7, 339: 7, 329: 7, 305: 7, 295: 7, 984: 6, 884: 6, 795: 6, 787:
 6, 743: 6, 722: 6, 718: 6, 681: 6, 654: 6, 652: 6, 629: 6, 623: 6, 606: 6, 588:
 6, 578: 6, 561: 6, 560: 6, 546: 6, 545: 6, 543: 6, 532: 6, 531: 6, 526: 6, 523:
 6, 513: 6, 512: 6, 485: 6, 483: 6, 479: 6, 470: 6, 454: 6, 449: 6, 447: 6, 443:
 6, 441: 6, 437: 6, 436: 6, 434: 6, 428: 6, 427: 6, 426: 6, 409: 6, 403: 6, 398:
 6, 390: 6, 387: 6, 378: 6, 367: 6, 358: 6, 353: 6, 348: 6, 332: 6, 321: 6, 306:
 6, 1309: 5, 1208: 5, 1173: 5, 972: 5, 929: 5, 927: 5, 871: 5, 863: 5, 852: 5,
 830: 5, 784: 5, 782: 5, 747: 5, 731: 5, 729: 5, 725: 5, 711: 5, 683: 5, 656: 5,
 655: 5, 642: 5, 636: 5, 624: 5, 611: 5, 609: 5, 603: 5, 600: 5, 585: 5, 579: 5,
 574: 5, 571: 5, 559: 5, 552: 5, 540: 5, 537: 5, 530: 5, 528: 5, 524: 5, 515: 5,
 506: 5, 504: 5, 500: 5, 493: 5, 491: 5, 480: 5, 478: 5, 477: 5, 473: 5, 472: 5,
 469: 5, 466: 5, 460: 5, 459: 5, 458: 5, 456: 5, 450: 5, 445: 5, 440: 5, 425: 5,
 423: 5, 420: 5, 414: 5, 413: 5, 383: 5, 365: 5, 343: 5, 335: 5, 330: 5, 323: 5,
 311: 5, 297: 5, 2069: 4, 1731: 4, 1499: 4, 1336: 4, 1146: 4, 1114: 4, 1109: 4,
 1101: 4, 1057: 4, 1045: 4, 1036: 4, 1025: 4, 981: 4, 963: 4, 961: 4, 943: 4,
 935: 4, 933: 4, 923: 4, 911: 4, 908: 4, 905: 4, 901: 4, 887: 4, 864: 4, 858: 4,
 842: 4, 837: 4, 836: 4, 826: 4, 808: 4, 791: 4, 781: 4, 772: 4, 762: 4, 733: 4,
 727: 4, 719: 4, 708: 4, 703: 4, 692: 4, 688: 4, 685: 4, 684: 4, 679: 4, 675: 4,
 674: 4, 673: 4, 672: 4, 668: 4, 663: 4, 658: 4, 653: 4, 651: 4, 650: 4, 647: 4,

641: 4, 635: 4, 621: 4, 618: 4, 614: 4, 613: 4, 601: 4, 592: 4, 590: 4, 589: 4,
 584: 4, 577: 4, 558: 4, 554: 4, 553: 4, 550: 4, 547: 4, 544: 4, 538: 4, 533: 4,
 522: 4, 511: 4, 505: 4, 501: 4, 496: 4, 495: 4, 489: 4, 486: 4, 465: 4, 463: 4,
 457: 4, 451: 4, 446: 4, 422: 4, 419: 4, 411: 4, 404: 4, 391: 4, 389: 4, 388: 4,
 386: 4, 382: 4, 366: 4, 359: 4, 344: 4, 341: 4, 331: 4, 294: 4, 3468: 3, 2891:
 3, 2714: 3, 2516: 3, 2286: 3, 2255: 3, 2158: 3, 2131: 3, 2087: 3, 1892: 3, 1883:
 3, 1878: 3, 1810: 3, 1779: 3, 1745: 3, 1708: 3, 1666: 3, 1637: 3, 1636: 3, 1635:
 3, 1629: 3, 1598: 3, 1589: 3, 1573: 3, 1526: 3, 1517: 3, 1469: 3, 1467: 3, 1435:
 3, 1399: 3, 1389: 3, 1367: 3, 1350: 3, 1325: 3, 1321: 3, 1314: 3, 1301: 3, 1299:
 3, 1291: 3, 1287: 3, 1285: 3, 1272: 3, 1256: 3, 1252: 3, 1223: 3, 1221: 3, 1217:
 3, 1207: 3, 1201: 3, 1196: 3, 1194: 3, 1186: 3, 1167: 3, 1158: 3, 1156: 3, 1155:
 3, 1143: 3, 1141: 3, 1135: 3, 1134: 3, 1111: 3, 1096: 3, 1083: 3, 1072: 3, 1068:
 3, 1049: 3, 1028: 3, 1009: 3, 992: 3, 990: 3, 986: 3, 983: 3, 975: 3, 974: 3,
 971: 3, 970: 3, 969: 3, 966: 3, 962: 3, 959: 3, 957: 3, 949: 3, 930: 3, 920: 3,
 912: 3, 909: 3, 906: 3, 902: 3, 894: 3, 892: 3, 883: 3, 882: 3, 879: 3, 877: 3,
 874: 3, 870: 3, 869: 3, 867: 3, 848: 3, 843: 3, 840: 3, 835: 3, 831: 3, 825: 3,
 821: 3, 818: 3, 812: 3, 805: 3, 804: 3, 800: 3, 799: 3, 797: 3, 789: 3, 783: 3,
 779: 3, 778: 3, 776: 3, 775: 3, 769: 3, 766: 3, 761: 3, 753: 3, 752: 3, 750: 3,
 748: 3, 744: 3, 737: 3, 735: 3, 723: 3, 721: 3, 720: 3, 717: 3, 714: 3, 707: 3,
 704: 3, 699: 3, 698: 3, 695: 3, 693: 3, 691: 3, 689: 3, 680: 3, 676: 3, 667: 3,
 666: 3, 662: 3, 661: 3, 660: 3, 657: 3, 645: 3, 639: 3, 638: 3, 637: 3, 632: 3,
 631: 3, 625: 3, 612: 3, 610: 3, 605: 3, 604: 3, 602: 3, 595: 3, 587: 3, 581: 3,
 576: 3, 575: 3, 570: 3, 569: 3, 564: 3, 562: 3, 556: 3, 542: 3, 520: 3, 519: 3,
 518: 3, 517: 3, 509: 3, 502: 3, 497: 3, 476: 3, 474: 3, 468: 3, 464: 3, 462: 3,
 455: 3, 448: 3, 444: 3, 415: 3, 401: 3, 399: 3, 379: 3, 307: 3, 12318: 2, 11221:
 2, 9750: 2, 7725: 2, 7645: 2, 6516: 2, 6421: 2, 6414: 2, 6224: 2, 5836: 2, 5767:
 2, 5670: 2, 5652: 2, 5445: 2, 4889: 2, 4835: 2, 4669: 2, 4217: 2, 4182: 2, 3988:
 2, 3774: 2, 3687: 2, 3651: 2, 3599: 2, 3457: 2, 3427: 2, 3411: 2, 3394: 2, 3388:
 2, 3383: 2, 3354: 2, 3349: 2, 3316: 2, 3298: 2, 3279: 2, 3270: 2, 3233: 2, 3179:
 2, 3177: 2, 3169: 2, 3089: 2, 3043: 2, 2985: 2, 2975: 2, 2952: 2, 2927: 2, 2884:
 2, 2879: 2, 2862: 2, 2843: 2, 2819: 2, 2787: 2, 2776: 2, 2756: 2, 2748: 2, 2708:
 2, 2685: 2, 2680: 2, 2639: 2, 2594: 2, 2590: 2, 2573: 2, 2568: 2, 2550: 2, 2489:
 2, 2450: 2, 2433: 2, 2414: 2, 2409: 2, 2384: 2, 2354: 2, 2338: 2, 2327: 2, 2311:
 2, 2300: 2, 2279: 2, 2260: 2, 2233: 2, 2228: 2, 2219: 2, 2214: 2, 2174: 2, 2166:
 2, 2161: 2, 2146: 2, 2136: 2, 2121: 2, 2118: 2, 2116: 2, 2073: 2, 2068: 2, 2064:
 2, 2054: 2, 2042: 2, 2039: 2, 2024: 2, 2003: 2, 1999: 2, 1997: 2, 1992: 2, 1981:
 2, 1976: 2, 1975: 2, 1969: 2, 1953: 2, 1947: 2, 1944: 2, 1935: 2, 1932: 2, 1928:
 2, 1923: 2, 1909: 2, 1897: 2, 1890: 2, 1880: 2, 1876: 2, 1874: 2, 1868: 2, 1859:
 2, 1845: 2, 1834: 2, 1832: 2, 1826: 2, 1813: 2, 1807: 2, 1798: 2, 1784: 2, 1777:
 2, 1776: 2, 1770: 2, 1766: 2, 1756: 2, 1755: 2, 1734: 2, 1728: 2, 1704: 2, 1682:
 2, 1674: 2, 1661: 2, 1654: 2, 1647: 2, 1645: 2, 1631: 2, 1630: 2, 1615: 2, 1606:
 2, 1594: 2, 1591: 2, 1584: 2, 1583: 2, 1580: 2, 1578: 2, 1569: 2, 1563: 2, 1556:
 2, 1555: 2, 1542: 2, 1530: 2, 1529: 2, 1525: 2, 1523: 2, 1520: 2, 1519: 2, 1512:
 2, 1508: 2, 1502: 2, 1501: 2, 1498: 2, 1497: 2, 1490: 2, 1489: 2, 1483: 2, 1481:
 2, 1480: 2, 1474: 2, 1472: 2, 1456: 2, 1449: 2, 1437: 2, 1429: 2, 1428: 2, 1423:
 2, 1417: 2, 1414: 2, 1413: 2, 1409: 2, 1408: 2, 1405: 2, 1400: 2, 1397: 2, 1396:
 2, 1395: 2, 1388: 2, 1387: 2, 1379: 2, 1370: 2, 1369: 2, 1355: 2, 1352: 2, 1346:
 2, 1342: 2, 1341: 2, 1340: 2, 1339: 2, 1335: 2, 1332: 2, 1329: 2, 1327: 2, 1322:

2, 1319: 2, 1312: 2, 1307: 2, 1306: 2, 1304: 2, 1295: 2, 1290: 2, 1289: 2, 1284:
 2, 1274: 2, 1268: 2, 1259: 2, 1258: 2, 1253: 2, 1250: 2, 1248: 2, 1247: 2, 1240:
 2, 1230: 2, 1229: 2, 1228: 2, 1227: 2, 1225: 2, 1219: 2, 1218: 2, 1213: 2, 1209:
 2, 1200: 2, 1199: 2, 1198: 2, 1197: 2, 1195: 2, 1193: 2, 1191: 2, 1188: 2, 1184:
 2, 1180: 2, 1172: 2, 1170: 2, 1169: 2, 1163: 2, 1162: 2, 1159: 2, 1154: 2, 1140:
 2, 1139: 2, 1136: 2, 1133: 2, 1132: 2, 1130: 2, 1128: 2, 1127: 2, 1126: 2, 1122:
 2, 1120: 2, 1118: 2, 1117: 2, 1116: 2, 1107: 2, 1106: 2, 1105: 2, 1097: 2, 1095:
 2, 1093: 2, 1091: 2, 1090: 2, 1086: 2, 1081: 2, 1080: 2, 1077: 2, 1074: 2, 1071:
 2, 1070: 2, 1065: 2, 1064: 2, 1063: 2, 1050: 2, 1048: 2, 1047: 2, 1044: 2, 1042:
 2, 1039: 2, 1037: 2, 1035: 2, 1033: 2, 1030: 2, 1014: 2, 1004: 2, 1003: 2, 1000:
 2, 999: 2, 965: 2, 956: 2, 951: 2, 950: 2, 944: 2, 942: 2, 940: 2, 936: 2, 932:
 2, 916: 2, 904: 2, 895: 2, 890: 2, 886: 2, 885: 2, 881: 2, 880: 2, 878: 2, 876:
 2, 866: 2, 862: 2, 861: 2, 854: 2, 847: 2, 845: 2, 844: 2, 839: 2, 834: 2, 829:
 2, 827: 2, 824: 2, 823: 2, 819: 2, 817: 2, 816: 2, 814: 2, 809: 2, 807: 2, 803:
 2, 798: 2, 793: 2, 788: 2, 785: 2, 774: 2, 773: 2, 760: 2, 759: 2, 758: 2, 757:
 2, 751: 2, 745: 2, 742: 2, 741: 2, 740: 2, 739: 2, 738: 2, 736: 2, 734: 2, 732:
 2, 730: 2, 726: 2, 724: 2, 715: 2, 713: 2, 709: 2, 705: 2, 701: 2, 697: 2, 694:
 2, 690: 2, 686: 2, 678: 2, 677: 2, 670: 2, 665: 2, 659: 2, 644: 2, 643: 2, 640:
 2, 630: 2, 628: 2, 627: 2, 620: 2, 619: 2, 617: 2, 615: 2, 598: 2, 596: 2, 594:
 2, 586: 2, 583: 2, 580: 2, 568: 2, 566: 2, 563: 2, 551: 2, 548: 2, 529: 2, 525:
 2, 521: 2, 514: 2, 510: 2, 503: 2, 499: 2, 498: 2, 492: 2, 475: 2, 150030: 1,
 119076: 1, 79358: 1, 66282: 1, 65306: 1, 65300: 1, 64997: 1, 62659: 1, 60999: 1,
 54288: 1, 53776: 1, 49333: 1, 48575: 1, 47310: 1, 45736: 1, 42919: 1, 42243: 1,
 42180: 1, 41961: 1, 41807: 1, 39682: 1, 38837: 1, 38520: 1, 38269: 1, 37755: 1,
 36906: 1, 36020: 1, 35706: 1, 35322: 1, 34886: 1, 34352: 1, 32952: 1, 32623: 1,
 32585: 1, 31576: 1, 31224: 1, 28994: 1, 27824: 1, 26398: 1, 25529: 1, 25393: 1,
 25360: 1, 25351: 1, 25333: 1, 24266: 1, 24045: 1, 23885: 1, 23651: 1, 23647: 1,
 23549: 1, 23266: 1, 23112: 1, 22064: 1, 21834: 1, 21600: 1, 21596: 1, 21591: 1,
 21162: 1, 21082: 1, 20475: 1, 20208: 1, 20042: 1, 20015: 1, 19912: 1, 19794: 1,
 19774: 1, 19135: 1, 19125: 1, 19064: 1, 18863: 1, 18628: 1, 18603: 1, 18480: 1,
 18443: 1, 18282: 1, 18222: 1, 18213: 1, 18077: 1, 18044: 1, 17953: 1, 17767: 1,
 17635: 1, 17499: 1, 17349: 1, 17290: 1, 17267: 1, 17084: 1, 17048: 1, 16961: 1,
 16901: 1, 16822: 1, 16815: 1, 16490: 1, 16482: 1, 16239: 1, 16220: 1, 15973: 1,
 15831: 1, 15710: 1, 15600: 1, 15507: 1, 15481: 1, 15405: 1, 15339: 1, 15298: 1,
 15271: 1, 15191: 1, 15141: 1, 15020: 1, 14845: 1, 14784: 1, 14761: 1, 14423: 1,
 14380: 1, 14313: 1, 14222: 1, 14192: 1, 14004: 1, 13964: 1, 13827: 1, 13825: 1,
 13745: 1, 13576: 1, 13493: 1, 13400: 1, 13232: 1, 13198: 1, 13116: 1, 13085: 1,
 13026: 1, 12860: 1, 12836: 1, 12819: 1, 12761: 1, 12746: 1, 12742: 1, 12660: 1,
 12650: 1, 12587: 1, 12530: 1, 12402: 1, 12371: 1, 12341: 1, 12242: 1, 12224: 1,
 12158: 1, 12107: 1, 12098: 1, 12095: 1, 12076: 1, 12073: 1, 12022: 1, 11936: 1,
 11924: 1, 11923: 1, 11910: 1, 11879: 1, 11793: 1, 11786: 1, 11763: 1, 11733: 1,
 11731: 1, 11725: 1, 11704: 1, 11697: 1, 11682: 1, 11569: 1, 11546: 1, 11533: 1,
 11358: 1, 11344: 1, 11294: 1, 11252: 1, 11162: 1, 11156: 1, 11106: 1, 11044: 1,
 10823: 1, 10816: 1, 10770: 1, 10738: 1, 10531: 1, 10523: 1, 10521: 1, 10506: 1,
 10298: 1, 10254: 1, 10116: 1, 10115: 1, 10106: 1, 10098: 1, 10092: 1, 9992: 1,
 9937: 1, 9919: 1, 9891: 1, 9880: 1, 9837: 1, 9816: 1, 9755: 1, 9672: 1, 9633: 1,
 9621: 1, 9568: 1, 9566: 1, 9542: 1, 9491: 1, 9431: 1, 9412: 1, 9363: 1, 9338: 1,
 9324: 1, 9272: 1, 9219: 1, 9179: 1, 9169: 1, 9127: 1, 9109: 1, 9067: 1, 9001: 1,

8981: 1, 8980: 1, 8962: 1, 8943: 1, 8915: 1, 8910: 1, 8884: 1, 8856: 1, 8847: 1,
8803: 1, 8707: 1, 8706: 1, 8694: 1, 8692: 1, 8664: 1, 8660: 1, 8637: 1, 8586: 1,
8548: 1, 8526: 1, 8514: 1, 8503: 1, 8493: 1, 8472: 1, 8462: 1, 8301: 1, 8261: 1,
8248: 1, 8150: 1, 8078: 1, 8074: 1, 8064: 1, 8053: 1, 8039: 1, 8031: 1, 8010: 1,
7998: 1, 7969: 1, 7961: 1, 7952: 1, 7926: 1, 7914: 1, 7891: 1, 7889: 1, 7879: 1,
7871: 1, 7868: 1, 7846: 1, 7838: 1, 7809: 1, 7804: 1, 7793: 1, 7756: 1, 7743: 1,
7731: 1, 7724: 1, 7662: 1, 7639: 1, 7607: 1, 7585: 1, 7562: 1, 7504: 1, 7480: 1,
7448: 1, 7411: 1, 7407: 1, 7404: 1, 7369: 1, 7344: 1, 7331: 1, 7306: 1, 7272: 1,
7224: 1, 7208: 1, 7206: 1, 7168: 1, 7128: 1, 7112: 1, 7109: 1, 7098: 1, 7076: 1,
7072: 1, 7071: 1, 7061: 1, 7052: 1, 7040: 1, 7034: 1, 6996: 1, 6985: 1, 6969: 1,
6954: 1, 6953: 1, 6949: 1, 6941: 1, 6934: 1, 6931: 1, 6918: 1, 6911: 1, 6857: 1,
6840: 1, 6830: 1, 6829: 1, 6815: 1, 6797: 1, 6792: 1, 6772: 1, 6768: 1, 6762: 1,
6748: 1, 6715: 1, 6707: 1, 6701: 1, 6694: 1, 6645: 1, 6643: 1, 6638: 1, 6634: 1,
6624: 1, 6566: 1, 6558: 1, 6489: 1, 6487: 1, 6483: 1, 6441: 1, 6376: 1, 6370: 1,
6316: 1, 6304: 1, 6292: 1, 6271: 1, 6270: 1, 6264: 1, 6237: 1, 6219: 1, 6201: 1,
6178: 1, 6149: 1, 6131: 1, 6123: 1, 6111: 1, 6103: 1, 6078: 1, 6057: 1, 6048: 1,
6044: 1, 6026: 1, 6021: 1, 6020: 1, 6012: 1, 5988: 1, 5975: 1, 5958: 1, 5931: 1,
5926: 1, 5892: 1, 5889: 1, 5838: 1, 5834: 1, 5833: 1, 5828: 1, 5795: 1, 5789: 1,
5782: 1, 5776: 1, 5771: 1, 5738: 1, 5733: 1, 5703: 1, 5700: 1, 5696: 1, 5689: 1,
5651: 1, 5639: 1, 5630: 1, 5623: 1, 5605: 1, 5592: 1, 5577: 1, 5576: 1, 5536: 1,
5527: 1, 5512: 1, 5505: 1, 5502: 1, 5501: 1, 5495: 1, 5480: 1, 5469: 1, 5461: 1,
5457: 1, 5454: 1, 5428: 1, 5425: 1, 5421: 1, 5419: 1, 5392: 1, 5375: 1, 5350: 1,
5348: 1, 5340: 1, 5339: 1, 5338: 1, 5333: 1, 5325: 1, 5319: 1, 5315: 1, 5308: 1,
5274: 1, 5273: 1, 5269: 1, 5251: 1, 5228: 1, 5208: 1, 5187: 1, 5162: 1, 5154: 1,
5127: 1, 5109: 1, 5093: 1, 5075: 1, 5071: 1, 5062: 1, 5051: 1, 5048: 1, 5037: 1,
5030: 1, 5027: 1, 5017: 1, 5008: 1, 5001: 1, 4996: 1, 4991: 1, 4990: 1, 4974: 1,
4947: 1, 4939: 1, 4923: 1, 4922: 1, 4909: 1, 4906: 1, 4898: 1, 4893: 1, 4876: 1,
4875: 1, 4871: 1, 4842: 1, 4838: 1, 4837: 1, 4833: 1, 4831: 1, 4822: 1, 4804: 1,
4790: 1, 4775: 1, 4773: 1, 4768: 1, 4756: 1, 4754: 1, 4749: 1, 4748: 1, 4743: 1,
4740: 1, 4731: 1, 4727: 1, 4701: 1, 4692: 1, 4685: 1, 4680: 1, 4679: 1, 4648: 1,
4645: 1, 4640: 1, 4618: 1, 4611: 1, 4598: 1, 4589: 1, 4583: 1, 4578: 1, 4563: 1,
4550: 1, 4538: 1, 4530: 1, 4527: 1, 4526: 1, 4509: 1, 4494: 1, 4491: 1, 4484: 1,
4482: 1, 4457: 1, 4447: 1, 4437: 1, 4432: 1, 4429: 1, 4427: 1, 4418: 1, 4416: 1,
4404: 1, 4395: 1, 4390: 1, 4388: 1, 4384: 1, 4380: 1, 4375: 1, 4372: 1, 4369: 1,
4368: 1, 4366: 1, 4361: 1, 4344: 1, 4342: 1, 4341: 1, 4340: 1, 4337: 1, 4336: 1,
4331: 1, 4330: 1, 4321: 1, 4317: 1, 4313: 1, 4311: 1, 4307: 1, 4301: 1, 4282: 1,
4279: 1, 4269: 1, 4263: 1, 4256: 1, 4253: 1, 4238: 1, 4222: 1, 4208: 1, 4192: 1,
4176: 1, 4175: 1, 4173: 1, 4148: 1, 4144: 1, 4141: 1, 4140: 1, 4139: 1, 4137: 1,
4134: 1, 4123: 1, 4110: 1, 4107: 1, 4094: 1, 4089: 1, 4087: 1, 4086: 1, 4081: 1,
4078: 1, 4074: 1, 4073: 1, 4059: 1, 4057: 1, 4053: 1, 4046: 1, 4037: 1, 4016: 1,
4012: 1, 3997: 1, 3991: 1, 3984: 1, 3982: 1, 3976: 1, 3975: 1, 3968: 1, 3961: 1,
3955: 1, 3948: 1, 3943: 1, 3938: 1, 3937: 1, 3935: 1, 3926: 1, 3923: 1, 3920: 1,
3915: 1, 3898: 1, 3897: 1, 3891: 1, 3884: 1, 3866: 1, 3851: 1, 3835: 1, 3834: 1,
3827: 1, 3816: 1, 3813: 1, 3812: 1, 3806: 1, 3802: 1, 3801: 1, 3796: 1, 3794: 1,
3788: 1, 3784: 1, 3782: 1, 3779: 1, 3766: 1, 3765: 1, 3764: 1, 3741: 1, 3739: 1,
3733: 1, 3727: 1, 3719: 1, 3718: 1, 3715: 1, 3707: 1, 3702: 1, 3692: 1, 3684: 1,
3682: 1, 3666: 1, 3659: 1, 3657: 1, 3644: 1, 3641: 1, 3640: 1, 3639: 1, 3638: 1,
3634: 1, 3633: 1, 3627: 1, 3621: 1, 3612: 1, 3610: 1, 3603: 1, 3602: 1, 3598: 1,

3591: 1, 3587: 1, 3586: 1, 3578: 1, 3573: 1, 3568: 1, 3563: 1, 3545: 1, 3541: 1,
3539: 1, 3536: 1, 3533: 1, 3532: 1, 3508: 1, 3506: 1, 3495: 1, 3489: 1, 3480: 1,
3474: 1, 3471: 1, 3465: 1, 3454: 1, 3452: 1, 3451: 1, 3447: 1, 3443: 1, 3438: 1,
3436: 1, 3434: 1, 3431: 1, 3426: 1, 3421: 1, 3413: 1, 3401: 1, 3396: 1, 3385: 1,
3374: 1, 3366: 1, 3365: 1, 3362: 1, 3361: 1, 3356: 1, 3355: 1, 3353: 1, 3351: 1,
3343: 1, 3327: 1, 3325: 1, 3319: 1, 3318: 1, 3314: 1, 3311: 1, 3307: 1, 3306: 1,
3302: 1, 3297: 1, 3292: 1, 3289: 1, 3288: 1, 3286: 1, 3285: 1, 3281: 1, 3271: 1,
3263: 1, 3262: 1, 3261: 1, 3260: 1, 3255: 1, 3254: 1, 3242: 1, 3238: 1, 3234: 1,
3228: 1, 3226: 1, 3224: 1, 3220: 1, 3219: 1, 3216: 1, 3211: 1, 3200: 1, 3199: 1,
3198: 1, 3197: 1, 3189: 1, 3167: 1, 3165: 1, 3161: 1, 3156: 1, 3153: 1, 3150: 1,
3143: 1, 3139: 1, 3137: 1, 3133: 1, 3132: 1, 3131: 1, 3128: 1, 3120: 1, 3107: 1,
3106: 1, 3104: 1, 3103: 1, 3100: 1, 3097: 1, 3094: 1, 3087: 1, 3075: 1, 3073: 1,
3072: 1, 3070: 1, 3062: 1, 3059: 1, 3057: 1, 3055: 1, 3050: 1, 3048: 1, 3041: 1,
3030: 1, 3029: 1, 3020: 1, 3013: 1, 3002: 1, 3001: 1, 2998: 1, 2997: 1, 2991: 1,
2986: 1, 2984: 1, 2982: 1, 2978: 1, 2968: 1, 2963: 1, 2960: 1, 2959: 1, 2955: 1,
2945: 1, 2938: 1, 2936: 1, 2933: 1, 2928: 1, 2926: 1, 2920: 1, 2914: 1, 2907: 1,
2905: 1, 2886: 1, 2885: 1, 2875: 1, 2872: 1, 2871: 1, 2870: 1, 2861: 1, 2856: 1,
2855: 1, 2848: 1, 2838: 1, 2831: 1, 2828: 1, 2827: 1, 2806: 1, 2804: 1, 2802: 1,
2791: 1, 2789: 1, 2783: 1, 2782: 1, 2768: 1, 2757: 1, 2750: 1, 2742: 1, 2741: 1,
2737: 1, 2733: 1, 2730: 1, 2725: 1, 2722: 1, 2715: 1, 2709: 1, 2700: 1, 2696: 1,
2694: 1, 2683: 1, 2672: 1, 2671: 1, 2667: 1, 2664: 1, 2654: 1, 2645: 1, 2641: 1,
2640: 1, 2638: 1, 2632: 1, 2631: 1, 2628: 1, 2627: 1, 2617: 1, 2616: 1, 2615: 1,
2614: 1, 2611: 1, 2610: 1, 2601: 1, 2600: 1, 2597: 1, 2595: 1, 2589: 1, 2586: 1,
2585: 1, 2584: 1, 2578: 1, 2576: 1, 2572: 1, 2569: 1, 2564: 1, 2563: 1, 2562: 1,
2556: 1, 2546: 1, 2539: 1, 2534: 1, 2532: 1, 2520: 1, 2519: 1, 2518: 1, 2517: 1,
2515: 1, 2512: 1, 2510: 1, 2509: 1, 2507: 1, 2506: 1, 2505: 1, 2504: 1, 2499: 1,
2494: 1, 2493: 1, 2487: 1, 2486: 1, 2485: 1, 2478: 1, 2477: 1, 2476: 1, 2473: 1,
2472: 1, 2466: 1, 2462: 1, 2461: 1, 2460: 1, 2455: 1, 2452: 1, 2437: 1, 2434: 1,
2429: 1, 2428: 1, 2427: 1, 2423: 1, 2420: 1, 2417: 1, 2416: 1, 2415: 1, 2413: 1,
2406: 1, 2401: 1, 2399: 1, 2397: 1, 2395: 1, 2392: 1, 2391: 1, 2390: 1, 2386: 1,
2380: 1, 2378: 1, 2368: 1, 2363: 1, 2362: 1, 2355: 1, 2353: 1, 2350: 1, 2346: 1,
2345: 1, 2344: 1, 2343: 1, 2340: 1, 2337: 1, 2336: 1, 2331: 1, 2325: 1, 2324: 1,
2322: 1, 2319: 1, 2317: 1, 2310: 1, 2304: 1, 2303: 1, 2301: 1, 2298: 1, 2292: 1,
2291: 1, 2284: 1, 2282: 1, 2280: 1, 2277: 1, 2274: 1, 2264: 1, 2261: 1, 2257: 1,
2253: 1, 2252: 1, 2250: 1, 2248: 1, 2247: 1, 2245: 1, 2241: 1, 2239: 1, 2238: 1,
2236: 1, 2230: 1, 2224: 1, 2217: 1, 2212: 1, 2205: 1, 2193: 1, 2192: 1, 2190: 1,
2189: 1, 2187: 1, 2186: 1, 2185: 1, 2184: 1, 2172: 1, 2165: 1, 2164: 1, 2162: 1,
2159: 1, 2157: 1, 2153: 1, 2151: 1, 2149: 1, 2144: 1, 2141: 1, 2138: 1, 2130: 1,
2125: 1, 2114: 1, 2112: 1, 2110: 1, 2109: 1, 2106: 1, 2105: 1, 2104: 1, 2101: 1,
2097: 1, 2096: 1, 2093: 1, 2090: 1, 2083: 1, 2080: 1, 2074: 1, 2060: 1, 2052: 1,
2050: 1, 2047: 1, 2036: 1, 2034: 1, 2033: 1, 2032: 1, 2027: 1, 2025: 1, 2017: 1,
2015: 1, 2014: 1, 2013: 1, 2012: 1, 2009: 1, 2008: 1, 2006: 1, 2004: 1, 2002: 1,
1998: 1, 1996: 1, 1993: 1, 1980: 1, 1978: 1, 1974: 1, 1972: 1, 1970: 1, 1967: 1,
1966: 1, 1960: 1, 1958: 1, 1957: 1, 1956: 1, 1954: 1, 1949: 1, 1942: 1, 1940: 1,
1937: 1, 1931: 1, 1930: 1, 1926: 1, 1925: 1, 1922: 1, 1921: 1, 1918: 1, 1913: 1,
1910: 1, 1904: 1, 1902: 1, 1898: 1, 1893: 1, 1889: 1, 1881: 1, 1879: 1, 1872: 1,
1870: 1, 1861: 1, 1855: 1, 1853: 1, 1852: 1, 1851: 1, 1847: 1, 1840: 1, 1837: 1,
1835: 1, 1830: 1, 1829: 1, 1828: 1, 1827: 1, 1825: 1, 1823: 1, 1822: 1, 1821: 1,

```

1820: 1, 1819: 1, 1818: 1, 1817: 1, 1815: 1, 1808: 1, 1805: 1, 1803: 1, 1802: 1,
1801: 1, 1795: 1, 1794: 1, 1793: 1, 1790: 1, 1788: 1, 1787: 1, 1781: 1, 1780: 1,
1778: 1, 1775: 1, 1774: 1, 1772: 1, 1771: 1, 1768: 1, 1765: 1, 1764: 1, 1762: 1,
1760: 1, 1758: 1, 1754: 1, 1753: 1, 1751: 1, 1748: 1, 1747: 1, 1741: 1, 1738: 1,
1737: 1, 1735: 1, 1733: 1, 1729: 1, 1727: 1, 1720: 1, 1717: 1, 1716: 1, 1713: 1,
1711: 1, 1710: 1, 1709: 1, 1705: 1, 1699: 1, 1697: 1, 1696: 1, 1694: 1, 1693: 1,
1692: 1, 1690: 1, 1687: 1, 1686: 1, 1685: 1, 1683: 1, 1679: 1, 1675: 1, 1671: 1,
1670: 1, 1669: 1, 1665: 1, 1664: 1, 1663: 1, 1660: 1, 1659: 1, 1657: 1, 1655: 1,
1653: 1, 1646: 1, 1643: 1, 1642: 1, 1640: 1, 1632: 1, 1626: 1, 1624: 1, 1623: 1,
1622: 1, 1620: 1, 1617: 1, 1614: 1, 1613: 1, 1612: 1, 1611: 1, 1608: 1, 1607: 1,
1605: 1, 1604: 1, 1603: 1, 1602: 1, 1600: 1, 1599: 1, 1595: 1, 1593: 1, 1590: 1,
1587: 1, 1586: 1, 1581: 1, 1579: 1, 1577: 1, 1574: 1, 1572: 1, 1570: 1, 1568: 1,
1566: 1, 1565: 1, 1564: 1, 1560: 1, 1559: 1, 1557: 1, 1554: 1, 1552: 1, 1548: 1,
1547: 1, 1546: 1, 1545: 1, 1541: 1, 1540: 1, 1539: 1, 1538: 1, 1534: 1, 1531: 1,
1528: 1, 1527: 1, 1516: 1, 1515: 1, 1509: 1, 1507: 1, 1505: 1, 1496: 1, 1495: 1,
1493: 1, 1488: 1, 1485: 1, 1484: 1, 1479: 1, 1478: 1, 1473: 1, 1471: 1, 1470: 1,
1468: 1, 1464: 1, 1463: 1, 1459: 1, 1457: 1, 1455: 1, 1452: 1, 1451: 1, 1445: 1,
1442: 1, 1441: 1, 1440: 1, 1439: 1, 1438: 1, 1430: 1, 1427: 1, 1426: 1, 1425: 1,
1418: 1, 1415: 1, 1411: 1, 1407: 1, 1406: 1, 1404: 1, 1403: 1, 1398: 1, 1392: 1,
1391: 1, 1386: 1, 1384: 1, 1382: 1, 1378: 1, 1377: 1, 1375: 1, 1372: 1, 1364: 1,
1361: 1, 1358: 1, 1357: 1, 1348: 1, 1347: 1, 1345: 1, 1344: 1, 1343: 1, 1338: 1,
1337: 1, 1334: 1, 1333: 1, 1331: 1, 1326: 1, 1323: 1, 1318: 1, 1317: 1, 1316: 1,
1313: 1, 1305: 1, 1298: 1, 1297: 1, 1294: 1, 1293: 1, 1288: 1, 1283: 1, 1282: 1,
1281: 1, 1276: 1, 1273: 1, 1271: 1, 1270: 1, 1269: 1, 1266: 1, 1264: 1, 1262: 1,
1261: 1, 1257: 1, 1255: 1, 1246: 1, 1244: 1, 1242: 1, 1235: 1, 1234: 1, 1232: 1,
1231: 1, 1222: 1, 1220: 1, 1216: 1, 1214: 1, 1212: 1, 1211: 1, 1206: 1, 1205: 1,
1203: 1, 1190: 1, 1183: 1, 1181: 1, 1178: 1, 1176: 1, 1175: 1, 1171: 1, 1165: 1,
1160: 1, 1157: 1, 1153: 1, 1152: 1, 1150: 1, 1149: 1, 1148: 1, 1147: 1, 1129: 1,
1125: 1, 1123: 1, 1115: 1, 1110: 1, 1108: 1, 1103: 1, 1102: 1, 1100: 1, 1099: 1,
1098: 1, 1094: 1, 1092: 1, 1089: 1, 1088: 1, 1085: 1, 1082: 1, 1076: 1, 1075: 1,
1069: 1, 1067: 1, 1061: 1, 1060: 1, 1059: 1, 1058: 1, 1055: 1, 1052: 1, 1046: 1,
1043: 1, 1041: 1, 1040: 1, 1034: 1, 1031: 1, 1029: 1, 1027: 1, 1023: 1, 1022: 1,
1021: 1, 1019: 1, 1018: 1, 1017: 1, 1013: 1, 1011: 1, 1008: 1, 1007: 1, 1006: 1,
1005: 1, 1002: 1, 1001: 1, 998: 1, 997: 1, 996: 1, 995: 1, 994: 1, 993: 1, 989:
1, 987: 1, 982: 1, 979: 1, 977: 1, 976: 1, 973: 1, 968: 1, 967: 1, 964: 1, 960:
1, 955: 1, 953: 1, 952: 1, 948: 1, 947: 1, 945: 1, 941: 1, 931: 1, 921: 1, 919:
1, 918: 1, 917: 1, 915: 1, 914: 1, 913: 1, 907: 1, 903: 1, 899: 1, 898: 1, 897:
1, 896: 1, 893: 1, 891: 1, 889: 1, 872: 1, 868: 1, 860: 1, 859: 1, 857: 1, 856:
1, 855: 1, 850: 1, 849: 1, 841: 1, 833: 1, 815: 1, 813: 1, 790: 1, 777: 1, 771:
1, 770: 1, 768: 1, 767: 1, 764: 1, 763: 1, 755: 1, 746: 1, 728: 1, 716: 1, 712:
1, 710: 1, 706: 1, 702: 1, 696: 1, 687: 1, 671: 1, 664: 1, 634: 1, 633: 1, 608:
1, 607: 1, 599: 1, 593: 1, 582: 1, 572: 1, 565: 1, 555: 1, 534: 1, 527: 1, 490:
1, 487: 1, 482: 1, 433: 1})

```

```

[73]: cv_log_error_array=[]
      for i in alpha:
          clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)

```

```

clf.fit(train_text_feature_onehotCoding, y_train)

sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_,
↪eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:",log_loss(y_cv,
↪predict_y, labels=clf.classes_, eps=1e-15))

```

```

For values of alpha = 1e-05 The log loss is: 1.2721829822908244
For values of alpha = 0.0001 The log loss is: 1.099474116965219
For values of alpha = 0.001 The log loss is: 1.1239105704089265
For values of alpha = 0.01 The log loss is: 1.2384219505980087
For values of alpha = 0.1 The log loss is: 1.4241586829998831
For values of alpha = 1 The log loss is: 1.642766769688418

```

```
[74]: print_best_alpha(alpha, cv_log_error_array)
```

```

The best alpha is: 0.0001
The best alpha index is: 1

```

```

[75]: # Simple model with only text data to evaluate its importance and check error
best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
↪random_state=42)
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
↪",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
↪log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
↪",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

For values of best alpha = 0.0001 The train log loss is: 0.648485394704729
For values of best alpha = 0.0001 The cross validation log loss is:
1.099474116965219
For values of best alpha = 0.0001 The test log loss is: 1.235792920444658

```

```

[76]: # Checking text overlap
def get_intersec_text(df):

```



```

df_text_vec = CountVectorizer(min_df=3)
df_text_fea = df_text_vec.fit_transform(df['TEXT'])
df_text_features = df_text_vec.get_feature_names()

df_text_fea_counts = df_text_fea.sum(axis=0).A1
df_text_fea_dict = dict(zip(list(df_text_features), df_text_fea_counts))
len1 = len(set(df_text_features))
len2 = len(set(train_text_features) & set(df_text_features))
return len1, len2

len1, len2 = get_intersec_text(test_df)
print(np.round((len2/len1)*100, 3), "% of word of test data appeared in train_
→data")
len1, len2 = get_intersec_text(cv_df)
print(np.round((len2/len1)*100, 3), "% of word of Cross Validation appeared in_
→train data")

```

96.274 % of word of test data appeared in train data
97.874 % of word of Cross Validation appeared in train data

0.8 Data prepration for Machine Learning models

```

[77]: # Functions delcaration

def report_log_loss(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    sig_clf_probs = sig_clf.predict_proba(test_x)
    return log_loss(test_y, sig_clf_probs, eps=1e-15)

# This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)

    A = (((C.T)/(C.sum(axis=1))).T)

    B = (C/C.sum(axis=0))
    labels = [1,2,3,4,5,6,7,8,9]
    # representing A in heatmap format
    print("-"*20, "Confusion matrix", "-"*20)
    plt.figure(figsize=(20,7))
    sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
→yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.show()

```

```

print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
→yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

# representing B in heatmap format
print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels,
→yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    pred_y = sig_clf.predict(test_x)

    # for calculating log_loss we will provide the array of probabilities
    →belongs to each class
    print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x)))
    # calculating the number of data points that are misclassified
    print("Number of mis-classified points :", np.count_nonzero((pred_y-
    →test_y))/test_y.shape[0])
    plot_confusion_matrix(test_y, pred_y)

# this function will be used just for naive bayes
# for the given indices, we will print the name of the features
# and we will check whether the feature present in the test point text or not
def get_impfeature_names(indices, text, gene, var, no_features):
    gene_count_vec = CountVectorizer()
    var_count_vec = CountVectorizer()
    text_count_vec = CountVectorizer(min_df=3)

    gene_vec = gene_count_vec.fit(train_df['Gene'])
    var_vec = var_count_vec.fit(train_df['Variation'])
    text_vec = text_count_vec.fit(train_df['TEXT'])

    fea1_len = len(gene_vec.get_feature_names())
    fea2_len = len(var_count_vec.get_feature_names())

```

```

word_present = 0
for i,v in enumerate(indices):
    if (v < fea1_len):
        word = gene_vec.get_feature_names()[v]
        yes_no = True if word == gene else False
        if yes_no:
            word_present += 1
            print(i, "Gene feature [{}] present in test data point [{}]"
        ↪format(word,yes_no))
    elif (v < fea1_len+fea2_len):
        word = var_vec.get_feature_names()[v-(fea1_len)]
        yes_no = True if word == var else False
        if yes_no:
            word_present += 1
            print(i, "variation feature [{}] present in test data point_
        ↪[{}]"
        ↪format(word,yes_no))
    else:
        word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
        yes_no = True if word in text.split() else False
        if yes_no:
            word_present += 1
            print(i, "Text feature [{}] present in test data point [{}]"
        ↪format(word,yes_no))

    print("Out of the top ",no_features," features ", word_present, "are_
    ↪present in query point")

```

Combining all 3 features together

```

[78]: # merging gene, variance and text features

# building train, test and cross validation data sets
# a = [[1, 2],
#      [3, 4]]
# b = [[4, 5],
#      [6, 7]]
# hstack(a, b) = [[1, 2, 4, 5],
#                 [ 3, 4, 6, 7]]

train_gene_var_onehotCoding =_
    ↪hstack((train_gene_feature_onehotCoding,train_variation_feature_onehotCoding))
test_gene_var_onehotCoding =_
    ↪hstack((test_gene_feature_onehotCoding,test_variation_feature_onehotCoding))
cv_gene_var_onehotCoding =_
    ↪hstack((cv_gene_feature_onehotCoding,cv_variation_feature_onehotCoding))

```

```

train_x_onehotCoding = hstack((train_gene_var_onehotCoding,
    ↳train_text_feature_onehotCoding)).tocsr()
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding,
    ↳test_text_feature_onehotCoding)).tocsr()
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding,
    ↳cv_text_feature_onehotCoding)).tocsr()
cv_y = np.array(list(cv_df['Class']))

train_gene_var_responseCoding = np.
    ↳hstack((train_gene_feature_responseCoding, train_variation_feature_responseCoding))
test_gene_var_responseCoding = np.
    ↳hstack((test_gene_feature_responseCoding, test_variation_feature_responseCoding))
cv_gene_var_responseCoding = np.
    ↳hstack((cv_gene_feature_responseCoding, cv_variation_feature_responseCoding))

train_x_responseCoding = np.hstack((train_gene_var_responseCoding,
    ↳train_text_feature_responseCoding))
test_x_responseCoding = np.hstack((test_gene_var_responseCoding,
    ↳test_text_feature_responseCoding))
cv_x_responseCoding = np.hstack((cv_gene_var_responseCoding,
    ↳cv_text_feature_responseCoding))

print("One hot encoding features :")
print("(number of data points * number of features) in train data = ",
    ↳train_x_onehotCoding.shape)
print("(number of data points * number of features) in test data = ",
    ↳test_x_onehotCoding.shape)
print("(number of data points * number of features) in cross validation data,
    ↳=", cv_x_onehotCoding.shape)
print(" Response encoding features :")
print("(number of data points * number of features) in train data = ",
    ↳train_x_responseCoding.shape)
print("(number of data points * number of features) in test data = ",
    ↳test_x_responseCoding.shape)
print("(number of data points * number of features) in cross validation data,
    ↳=", cv_x_responseCoding.shape)

```

One hot encoding features :

(number of data points * number of features) in train data = (2124, 55457)

(number of data points * number of features) in test data = (665, 55457)

(number of data points * number of features) in cross validation data = (532, 55457)

Response encoding features :

(number of data points * number of features) in train data = (2124, 27)

(number of data points * number of features) in test data = (665, 27)

(number of data points * number of features) in cross validation data = (532, 27)

0.9 Building Machine Learning models

0.10 Naive Bayes

```
[79]: # http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.  
      ↪ MultinomialNB.html  
alpha = [0.00001, 0.0001, 0.001, 0.1, 1, 10, 100, 1000]  
cv_log_error_array = []  
for i in alpha:  
    print("for alpha =", i)  
    # MultinomialNB is used for multi class classification  
    clf = MultinomialNB(alpha=i)  
    clf.fit(train_x_onehotCoding, train_y)  
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")  
    sig_clf.fit(train_x_onehotCoding, train_y)  
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)  
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.  
    ↪ classes_, eps=1e-15))  
    # to avoid rounding error while multiplying probabilities we use  
    ↪ log-probability estimates  
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))  
  
print_best_alpha(alpha, cv_log_error_array)
```

```
for alpha = 1e-05  
Log Loss : 1.2772367409831071  
for alpha = 0.0001  
Log Loss : 1.2757405184068615  
for alpha = 0.001  
Log Loss : 1.270251486108716  
for alpha = 0.1  
Log Loss : 1.2464245671312086  
for alpha = 1  
Log Loss : 1.256857498516061  
for alpha = 10  
Log Loss : 1.3833717634434648  
for alpha = 100  
Log Loss : 1.3600948522454712  
for alpha = 1000  
Log Loss : 1.309119867422757
```

The best alpha is: 0.1
The best alpha index is: 3

```
[80]: best_alpha = np.argmin(cv_log_error_array)
      clf = MultinomialNB(alpha=alpha[best_alpha])
      clf.fit(train_x_onehotCoding, train_y)
      sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
      sig_clf.fit(train_x_onehotCoding, train_y)

      predict_y = sig_clf.predict_proba(train_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
      ↪", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
      predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The cross validation_
      ↪log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
      predict_y = sig_clf.predict_proba(test_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
      ↪", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```


For values of best alpha = 0.1 The train log loss is: 0.8526757766833339
For values of best alpha = 0.1 The cross validation log loss is:
1.2464245671312086
For values of best alpha = 0.1 The test log loss is: 1.3128752779765938

```
[81]: clf = MultinomialNB(alpha=alpha[best_alpha])
      clf.fit(train_x_onehotCoding, train_y)
      sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
      sig_clf.fit(train_x_onehotCoding, train_y)
      sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
      # to avoid rounding error while multiplying probabilities we use log-probability_
      ↪estimates
      print("Log Loss :", log_loss(cv_y, sig_clf_probs))
      print("Number of misclassified point :", np.count_nonzero((sig_clf.
      ↪predict(cv_x_onehotCoding) - cv_y))/cv_y.shape[0])
      plot_confusion_matrix(cv_y, sig_clf.predict(cv_x_onehotCoding.toarray()))
```


Log Loss : 1.2464245671312086
Number of misclassified point : 0.37218045112781956
----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



From precision and recall matrix diagonal values show correctly predicted values Checking these matrixes we can also see between which classes is confusion happening

Interpretability of Naive Bayes

```
[82]: # checked item is defined apriory, without any meaning
# we can check whichever item we want
# defining item to be checked
test_point_index = 2
# important features to be printed
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.
    ↳predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT']).
    ↳iloc[test_point_index],test_df['Gene'].
    ↳iloc[test_point_index],test_df['Variation'].iloc[test_point_index],_
    ↳no_feature)
```

Predicted Class : 7

Predicted Class Probabilities: [[0.077 0.0634 0.014 0.1058 0.0332 0.0317
0.6676 0.004 0.0033]]

Actual Class : 7

```
-----
16 Text feature [kinase] present in test data point [True]
17 Text feature [presence] present in test data point [True]
18 Text feature [activating] present in test data point [True]
19 Text feature [downstream] present in test data point [True]
```


20 Text feature [independent] present in test data point [True]
25 Text feature [inhibitor] present in test data point [True]
26 Text feature [well] present in test data point [True]
27 Text feature [expressing] present in test data point [True]
28 Text feature [recently] present in test data point [True]
29 Text feature [previously] present in test data point [True]
30 Text feature [activation] present in test data point [True]
31 Text feature [cell] present in test data point [True]
32 Text feature [contrast] present in test data point [True]
33 Text feature [cells] present in test data point [True]
34 Text feature [also] present in test data point [True]
35 Text feature [potential] present in test data point [True]
36 Text feature [shown] present in test data point [True]
37 Text feature [showed] present in test data point [True]
38 Text feature [growth] present in test data point [True]
39 Text feature [found] present in test data point [True]
40 Text feature [10] present in test data point [True]
41 Text feature [however] present in test data point [True]
42 Text feature [suggest] present in test data point [True]
43 Text feature [compared] present in test data point [True]
44 Text feature [similar] present in test data point [True]
45 Text feature [higher] present in test data point [True]
46 Text feature [observed] present in test data point [True]
47 Text feature [addition] present in test data point [True]
48 Text feature [mutations] present in test data point [True]
49 Text feature [treated] present in test data point [True]
50 Text feature [described] present in test data point [True]
52 Text feature [may] present in test data point [True]
53 Text feature [obtained] present in test data point [True]
54 Text feature [factor] present in test data point [True]
55 Text feature [proliferation] present in test data point [True]
57 Text feature [1a] present in test data point [True]
58 Text feature [approximately] present in test data point [True]
59 Text feature [interestingly] present in test data point [True]
60 Text feature [inhibition] present in test data point [True]
61 Text feature [using] present in test data point [True]
62 Text feature [3a] present in test data point [True]
63 Text feature [respectively] present in test data point [True]
64 Text feature [identified] present in test data point [True]
65 Text feature [activated] present in test data point [True]
66 Text feature [3b] present in test data point [True]
67 Text feature [reported] present in test data point [True]
68 Text feature [including] present in test data point [True]
69 Text feature [12] present in test data point [True]
70 Text feature [phosphorylation] present in test data point [True]
71 Text feature [studies] present in test data point [True]
72 Text feature [followed] present in test data point [True]
73 Text feature [various] present in test data point [True]

```

74 Text feature [total] present in test data point [True]
75 Text feature [inhibitors] present in test data point [True]
76 Text feature [new] present in test data point [True]
77 Text feature [three] present in test data point [True]
78 Text feature [mechanism] present in test data point [True]
79 Text feature [fig] present in test data point [True]
81 Text feature [confirmed] present in test data point [True]
82 Text feature [occur] present in test data point [True]
83 Text feature [either] present in test data point [True]
84 Text feature [molecular] present in test data point [True]
85 Text feature [consistent] present in test data point [True]
87 Text feature [confirm] present in test data point [True]
88 Text feature [sensitive] present in test data point [True]
89 Text feature [suggests] present in test data point [True]
90 Text feature [although] present in test data point [True]
91 Text feature [two] present in test data point [True]
92 Text feature [mutation] present in test data point [True]
93 Text feature [different] present in test data point [True]
94 Text feature [due] present in test data point [True]
95 Text feature [increase] present in test data point [True]
96 Text feature [figure] present in test data point [True]
97 Text feature [show] present in test data point [True]
98 Text feature [revealed] present in test data point [True]
99 Text feature [enhanced] present in test data point [True]
Out of the top 100 features 76 are present in query point

```

May have to check comparing with a Naive Bayes builded using only text data as input

Looking at another item

```

[83]: test_point_index = 100
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.
    ↳predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].
    ↳iloc[test_point_index],test_df['Gene'].
    ↳iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
    ↳no_feature)

```

Predicted Class : 5

Predicted Class Probabilities: [[0.1016 0.0834 0.0185 0.1394 0.4425 0.0423
0.1628 0.0052 0.0043]]

Actual Class : 4

5 Text feature [neutral] present in test data point [True]
8 Text feature [assays] present in test data point [True]
12 Text feature [functional] present in test data point [True]
Out of the top 100 features 3 are present in query point

1 K Nearest Neighbour Classification

```
[84]: alpha = [5, 11, 15, 21, 31, 41, 51, 99]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = KNeighborsClassifier(n_neighbors=i)
    clf.fit(train_x_responseCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_responseCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.
    ↪classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use ↪
    ↪log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))
print("")
print_best_alpha(alpha, cv_log_error_array)
```

```
for alpha = 5
Log Loss : 1.0587095490258613
for alpha = 11
Log Loss : 1.074858782876268
for alpha = 15
Log Loss : 1.0830905463973681
for alpha = 21
Log Loss : 1.095072992000385
for alpha = 31
Log Loss : 1.114826777681961
for alpha = 41
Log Loss : 1.13301441582936
for alpha = 51
Log Loss : 1.1478034524781693
for alpha = 99
Log Loss : 1.156847776021154
```

The best alpha is: 5

The best alpha index is: 0

May want to dense alpha values between 5 and 15 for a better alpha, but for the purpose of this kernel will take 11

```
[85]: best_alpha = np.argmin(cv_log_error_array)
      clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
      clf.fit(train_x_responseCoding, train_y)
      sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
      sig_clf.fit(train_x_responseCoding, train_y)


      predict_y = sig_clf.predict_proba(train_x_responseCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
            ↪", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
      predict_y = sig_clf.predict_proba(cv_x_responseCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The cross validation_
            ↪log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
      predict_y = sig_clf.predict_proba(test_x_responseCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
            ↪", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

For values of best alpha = 5 The train log loss is: 0.44790408188976055
 For values of best alpha = 5 The cross validation log loss is:
 1.0587095490258613
 For values of best alpha = 5 The test log loss is: 1.1368923111769251

```
[86]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
      predict_and_plot_confusion_matrix(train_x_responseCoding, train_y,
            ↪cv_x_responseCoding, cv_y, clf)
```

Log loss : 1.0587095490258613
 Number of mis-classified points : 0.35902255639097747

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



```
[87]: # Lets look at few test points
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 1
predicted_cls = sig_clf.predict(test_x_responseCoding[0].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1,-1), alpha[best_alpha])
```

```
print("The ",alpha[best_alpha]," nearest neighbours of the test points belongs_
↳to classes",train_y[neighbors[1][0]])
print("Fequency of nearest points :",Counter(train_y[neighbors[1][0]))
```

Predicted Class : 1

Actual Class : 7

The 5 nearest neighbours of the test points belongs to classes [7 7 7 7 7]

Fequency of nearest points : Counter({7: 5})

```
[88]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 100

predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].
↳reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1,
↳-1), alpha[best_alpha])
print("the k value for knn is",alpha[best_alpha],"and the nearest neighbours of_
↳the test points belongs to classes",train_y[neighbors[1][0]])
print("Fequency of nearest points :",Counter(train_y[neighbors[1][0]))
```

Predicted Class : 4

Actual Class : 4

the k value for knn is 5 and the nearest neighbours of the test points belongs to classes [5 4 4 4 4]

Fequency of nearest points : Counter({4: 4, 5: 1})

2 Logistic Regression

logistic regression with ballancing of classes

```
[90]: alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2',
↳loss='log', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
```

```

        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.
↪classes_, eps=1e-15))
        # to avoid rounding error while multiplying probabilities we use
↪log-probability estimates
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))
print("")
print_best_alpha(alpha, cv_log_error_array)

```

```

for alpha = 1e-06
Log Loss : 1.3497093406473863
for alpha = 1e-05
Log Loss : 1.2950280592437216
for alpha = 0.0001
Log Loss : 1.1110092263689428
for alpha = 0.001
Log Loss : 1.074115673621701
for alpha = 0.01
Log Loss : 1.1314175972482177
for alpha = 0.1
Log Loss : 1.4648396344774646
for alpha = 1
Log Loss : 1.694290394268692
for alpha = 10
Log Loss : 1.7223210557367996
for alpha = 100
Log Loss : 1.7253925391967233

```

```

The best alpha is: 0.001
The best alpha index is: 3

```

```

[91]: best_alpha = np.argmin(cv_log_error_array)
      clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha],
↪penalty='l2', loss='log', random_state=42)
      clf.fit(train_x_onehotCoding, train_y)
      sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
      sig_clf.fit(train_x_onehotCoding, train_y)

      predict_y = sig_clf.predict_proba(train_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
↪",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
      predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The cross validation
↪log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
      predict_y = sig_clf.predict_proba(test_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
↪",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```


For values of best alpha = 0.001 The train log loss is: 0.5105776633624867
 For values of best alpha = 0.001 The cross validation log loss is:
 1.074115673621701
 For values of best alpha = 0.001 The test log loss is: 1.1526170372094005

```
[92]: clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha],
    ↪penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,
    ↪cv_x_onehotCoding, cv_y, clf)
```


Log loss : 1.074115673621701

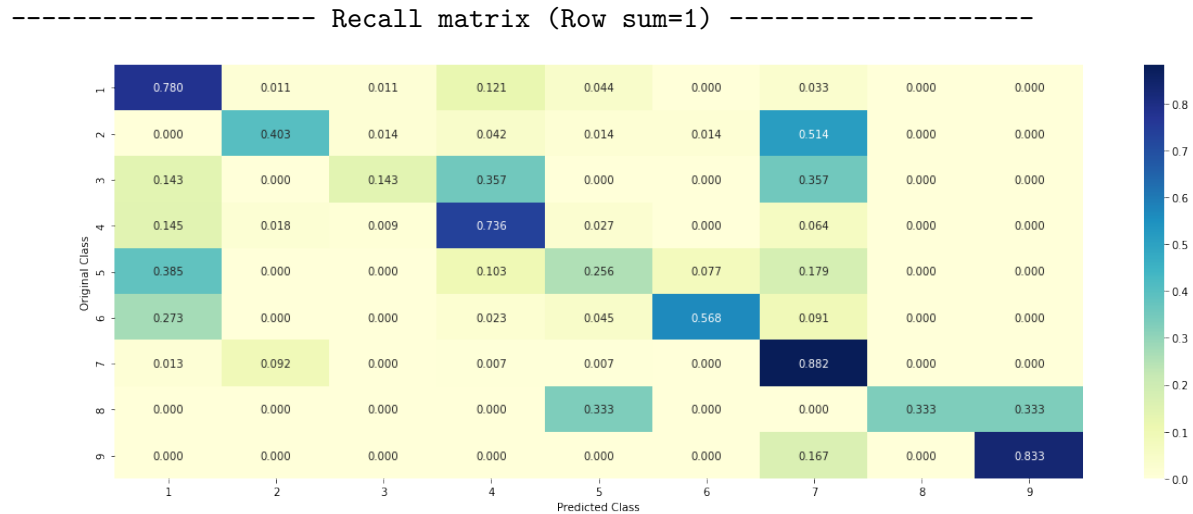
Number of mis-classified points : 0.325187969924812

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----





Logistic regression interpreability

```
[93]: def get_imp_feature_names(text, indices, removed_ind = []):
    word_present = 0
    tabulte_list = []
    incresingorder_ind = 0
    for i in indices:
        if i < train_gene_feature_onehotCoding.shape[1]:
            tabulte_list.append([incresingorder_ind, "Gene", "Yes"])
        elif i < 18:
            tabulte_list.append([incresingorder_ind, "Variation", "Yes"])
        if ((i > 17) & (i not in removed_ind)) :
            word = train_text_features[i]
            yes_no = True if word in text.split() else False
            if yes_no:
                word_present += 1
            tabulte_list.append([incresingorder_ind, train_text_features[i],
↪yes_no])
            incresingorder_ind += 1
        print(word_present, "most important features are present in our query",
↪point")
        print("-"*50)
        print("The features that are most important of the ", predicted_cls[0], "
↪class:")
        print(tabulate(tabulte_list, headers=["Index", "Feature name", "Present or
↪Not"]))
```

```
[94]: # from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha],
↪penalty='l2', loss='log', random_state=42)
```

```

clf.fit(train_x_onehotCoding,train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.
    ↳predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].
    ↳iloc[test_point_index],test_df['Gene'].
    ↳iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
    ↳no_feature)

```

```

Predicted Class : 7
Predicted Class Probabilities: [[0.0831 0.2388 0.0223 0.0914 0.0474 0.0592
0.4401 0.0075 0.0103]]
Actual Class : 7

```

```

-----
24 Text feature [constitutively] present in test data point [True]
35 Text feature [nude] present in test data point [True]
45 Text feature [activated] present in test data point [True]
49 Text feature [constitutive] present in test data point [True]
79 Text feature [oncogene] present in test data point [True]
137 Text feature [extracellular] present in test data point [True]
143 Text feature [tk1] present in test data point [True]
179 Text feature [activation] present in test data point [True]
211 Text feature [ligand] present in test data point [True]
239 Text feature [stably] present in test data point [True]
242 Text feature [tk2] present in test data point [True]
272 Text feature [transformation] present in test data point [True]
289 Text feature [kinase] present in test data point [True]
328 Text feature [tyrosine] present in test data point [True]
349 Text feature [ba] present in test data point [True]
366 Text feature [murine] present in test data point [True]
400 Text feature [f3] present in test data point [True]
Out of the top 500 features 17 are present in query point

```

```

[95]: test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.
    ↳predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]

```

```

print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].
    ↳iloc[test_point_index],test_df['Gene'].
    ↳iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
    ↳no_feature)

```

Predicted Class : 4

Predicted Class Probabilities: [[0.2025 0.0121 0.0157 0.3599 0.308 0.0915
0.0034 0.004 0.0029]]

Actual Class : 4

```

-----
137 Text feature [instability] present in test data point [True]
190 Text feature [382] present in test data point [True]
202 Text feature [adapted] present in test data point [True]
311 Text feature [cmmrd] present in test data point [True]
468 Text feature [microsatellite] present in test data point [True]
Out of the top 500 features 5 are present in query point

```

Notes from some different runs: None of top 100 and 2 of top 500 top features on this particular missclassified point
None of top 100 and 18 of top 500 top features on this particular missclassified point
None of top 200 and 7 of top 500 top features on this particular missclassified point
None of top 100 seems constant confusion (logically)

Logistic regression without balancing

```

[96]: alpha = [10 ** x for x in range(-6, 1)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.
    ↳classes_, eps=1e-15))
    print("Log Loss :",log_loss(cv_y, sig_clf_probs))
print("")
print_best_alpha(alpha, cv_log_error_array)

```

```

for alpha = 1e-06
Log Loss : 1.2908779040202831
for alpha = 1e-05
Log Loss : 1.2590311356289654
for alpha = 0.0001
Log Loss : 1.1075231754795696
for alpha = 0.001
Log Loss : 1.0815711224720403
for alpha = 0.01

```

```
Log Loss : 1.1773799272301868
for alpha = 0.1
Log Loss : 1.3256253560462177
for alpha = 1
Log Loss : 1.5740467277372516
```

```
The best alpha is: 0.001
The best alpha index is: 3
```


```
[97]: best_alpha = np.argmin(cv_log_error_array)
      clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
      ↪random_state=42)
      clf.fit(train_x_onehotCoding, train_y)
      sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
      sig_clf.fit(train_x_onehotCoding, train_y)

      predict_y = sig_clf.predict_proba(train_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
      ↪", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
      predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The cross validation
      ↪log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
      predict_y = sig_clf.predict_proba(test_x_onehotCoding)
      print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
      ↪", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```


```
For values of best alpha = 0.001 The train log loss is: 0.5085015257966655
For values of best alpha = 0.001 The cross validation log loss is:
1.0815711224720403
For values of best alpha = 0.001 The test log loss is: 1.1634584797518273
```

```
[98]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
      ↪random_state=42)
      predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,
      ↪cv_x_onehotCoding, cv_y, clf)
```


```
Log loss : 1.0815711224720403
Number of mis-classified points : 0.32706766917293234
----- Confusion matrix -----
```



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



Testing query point and interpretability

```
[99]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
    ↪random_state=42)
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.
    ↪predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].
    ↪iloc[test_point_index],test_df['Gene'].
    ↪iloc[test_point_index],test_df['Variation'].iloc[test_point_index],
    ↪no_feature)
```

Predicted Class : 7

Predicted Class Probabilities: [[0.0738 0.2264 0.0285 0.079 0.0496 0.0592
0.4589 0.0096 0.015]]

Actual Class : 7

```
-----
52 Text feature [constitutively] present in test data point [True]
110 Text feature [nude] present in test data point [True]
122 Text feature [activated] present in test data point [True]
125 Text feature [constitutive] present in test data point [True]
180 Text feature [oncogene] present in test data point [True]
273 Text feature [stably] present in test data point [True]
278 Text feature [extracellular] present in test data point [True]
289 Text feature [tk1] present in test data point [True]
```

```

297 Text feature [activation] present in test data point [True]
317 Text feature [transformation] present in test data point [True]
339 Text feature [ba] present in test data point [True]
341 Text feature [kinase] present in test data point [True]
396 Text feature [f3] present in test data point [True]
425 Text feature [ligand] present in test data point [True]
447 Text feature [tk2] present in test data point [True]
486 Text feature [tyrosine] present in test data point [True]
Out of the top 500 features 16 are present in query point

```

3 Linear Support Vector Machines

```

[100]: alpha = [10 ** x for x in range(-5, 3)]
cv_log_error_array = []
for i in alpha:
    print("for C =", i)
#     clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2',
    ↪ loss='hinge', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.
    ↪ classes_, eps=1e-15))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))
print("")
print_best_alpha(alpha, cv_log_error_array)

```

```

for C = 1e-05
Log Loss : 1.332081971282043
for C = 0.0001
Log Loss : 1.241496544243648
for C = 0.001
Log Loss : 1.0951933226853459
for C = 0.01
Log Loss : 1.081277561931896
for C = 0.1
Log Loss : 1.3763349868914128
for C = 1
Log Loss : 1.7122342245702478
for C = 10
Log Loss : 1.7259457630017707
for C = 100
Log Loss : 1.725945566813617

```

The best alpha is: 0.01

The best alpha index is: 3

```
[101]: best_alpha = np.argmin(cv_log_error_array)
# clf = SVC(C=i,kernel='linear',probability=True, class_weight='balanced')
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha],
    ↪penalty='l2', loss='hinge', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:
    ↪",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation
    ↪log loss is:",log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:
    ↪",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
```

For values of best alpha = 0.01 The train log loss is: 0.7177803857393527

For values of best alpha = 0.01 The cross validation log loss is:

1.081277561931896

For values of best alpha = 0.01 The test log loss is: 1.19027652332489

```
[102]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge',
    ↪random_state=42,class_weight='balanced')
predict_and_plot_confusion_matrix(train_x_onehotCoding,
    ↪train_y,cv_x_onehotCoding,cv_y, clf)
```

Log loss : 1.081277561931896

Number of mis-classified points : 0.34022556390977443


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



```
[103]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge',
    ↪ random_state=42)
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 1
# test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
```

```

print("Predicted Class Probabilities:", np.round(sig_clf.
    ↳predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].
    ↳iloc[test_point_index],test_df['Gene'].
    ↳iloc[test_point_index],test_df['Variation'].iloc[test_point_index],↳
    ↳no_feature)

```

Predicted Class : 7

Predicted Class Probabilities: [[0.1217 0.2478 0.0217 0.1349 0.0571 0.0543
0.3469 0.0075 0.008]]

Actual Class : 7

```

-----
36 Text feature [constitutively] present in test data point [True]
38 Text feature [nude] present in test data point [True]
74 Text feature [activated] present in test data point [True]
135 Text feature [tk1] present in test data point [True]
159 Text feature [oncogene] present in test data point [True]
166 Text feature [stably] present in test data point [True]
182 Text feature [constitutive] present in test data point [True]
223 Text feature [tk2] present in test data point [True]
239 Text feature [extracellular] present in test data point [True]
306 Text feature [activation] present in test data point [True]
398 Text feature [grew] present in test data point [True]
469 Text feature [kinase] present in test data point [True]
Out of the top 500 features 12 are present in query point

```

4 Random Forest Classifier

```

[104]: alpha = [100,200,500,1000,2000]
max_depth = [5, 10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini',↳
    ↳max_depth=j, random_state=42, n_jobs=-1)
        clf.fit(train_x_onehotCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_onehotCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.
    ↳classes_, eps=1e-15))
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))

```

```

for n_estimators = 100 and max depth = 5
Log Loss : 1.2565422522238885
for n_estimators = 100 and max depth = 10
Log Loss : 1.164423370599573
for n_estimators = 200 and max depth = 5
Log Loss : 1.2479314765592489
for n_estimators = 200 and max depth = 10
Log Loss : 1.158147703478899
for n_estimators = 500 and max depth = 5
Log Loss : 1.2401119952488662
for n_estimators = 500 and max depth = 10
Log Loss : 1.1520341944684371
for n_estimators = 1000 and max depth = 5
Log Loss : 1.2407551791690425
for n_estimators = 1000 and max depth = 10
Log Loss : 1.1517907344797522
for n_estimators = 2000 and max depth = 5
Log Loss : 1.2374495389483529
for n_estimators = 2000 and max depth = 10
Log Loss : 1.1513661957058137

```

```

[105]: best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)],
    ↳ criterion='gini', max_depth=max_depth[int(best_alpha%2)], random_state=42,
    ↳ n_jobs=-1)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The train_
    ↳ log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The cross_
    ↳ validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_,
    ↳ eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The test_
    ↳ log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

For values of best estimator = 2000 The train log loss is: 0.694672023605049
For values of best estimator = 2000 The cross validation log loss is:
1.1513661957058137
For values of best estimator = 2000 The test log loss is: 1.1785550523354904

```

[106]:

```


clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)],
    ↪criterion='gini', max_depth=max_depth[int(best_alpha%2)], random_state=42,
    ↪n_jobs=-1)
predict_and_plot_confusion_matrix(train_x_onehotCoding,
    ↪train_y,cv_x_onehotCoding,cv_y, clf)

```


Log loss : 1.1513661957058137

Number of mis-classified points : 0.38345864661654133


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



```
[107]: # test_point_index = 10
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)],
    ↳ criterion='gini', max_depth=max_depth[int(best_alpha%2)], random_state=42,
    ↳ n_jobs=-1)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.
    ↳ predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
get_impfeature_names(indices[:no_feature], test_df['TEXT'].
    ↳ iloc[test_point_index], test_df['Gene'].
    ↳ iloc[test_point_index], test_df['Variation'].iloc[test_point_index],
    ↳ no_feature)
```

```
Predicted Class : 7
Predicted Class Probabilities: [[0.1098 0.139 0.0235 0.1534 0.0586 0.0497
0.4515 0.0071 0.0075]]
Actual Class : 7
```

```
-----
0 Text feature [kinase] present in test data point [True]
2 Text feature [phosphorylation] present in test data point [True]
3 Text feature [tyrosine] present in test data point [True]
4 Text feature [activated] present in test data point [True]
```

```

6 Text feature [activation] present in test data point [True]
10 Text feature [constitutive] present in test data point [True]
11 Text feature [function] present in test data point [True]
17 Text feature [cells] present in test data point [True]
19 Text feature [oncogenic] present in test data point [True]
21 Text feature [constitutively] present in test data point [True]
23 Text feature [receptor] present in test data point [True]
35 Text feature [ba] present in test data point [True]
42 Text feature [cell] present in test data point [True]
45 Text feature [proliferation] present in test data point [True]
46 Text feature [f3] present in test data point [True]
53 Text feature [oncogene] present in test data point [True]
62 Text feature [expression] present in test data point [True]
65 Text feature [lines] present in test data point [True]
67 Text feature [extracellular] present in test data point [True]
76 Text feature [transformation] present in test data point [True]
86 Text feature [proteins] present in test data point [True]
95 Text feature [il] present in test data point [True]
96 Text feature [dna] present in test data point [True]
98 Text feature [presence] present in test data point [True]
99 Text feature [mutant] present in test data point [True]
Out of the top 100 features 25 are present in query point

```

4.1 RF with Response Coding

```

[108]: alpha = [10,50,100,200,500,1000]
max_depth = [2,3,5,10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators = ", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini',
→max_depth=j, random_state=42, n_jobs=-1)
        clf.fit(train_x_responseCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_responseCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.
→classes_, eps=1e-15))
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)],
→criterion='gini', max_depth=max_depth[int(best_alpha%4)], random_state=42,
→n_jobs=-1)
clf.fit(train_x_responseCoding, train_y)

```

```

sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The train log_
↪loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The cross_
↪validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_,
↪eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The test log_
↪loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for n_estimators = 10 and max depth = 2
Log Loss : 2.019011558757825
for n_estimators = 10 and max depth = 3
Log Loss : 1.6836830088809915
for n_estimators = 10 and max depth = 5
Log Loss : 1.4709967015832548
for n_estimators = 10 and max depth = 10
Log Loss : 1.9066898602863065
for n_estimators = 50 and max depth = 2
Log Loss : 1.5898458702429565
for n_estimators = 50 and max depth = 3
Log Loss : 1.3833003257090064
for n_estimators = 50 and max depth = 5
Log Loss : 1.2833788362382992
for n_estimators = 50 and max depth = 10
Log Loss : 1.7657298571315907
for n_estimators = 100 and max depth = 2
Log Loss : 1.4601838920233587
for n_estimators = 100 and max depth = 3
Log Loss : 1.40911038483864
for n_estimators = 100 and max depth = 5
Log Loss : 1.3176183811501851
for n_estimators = 100 and max depth = 10
Log Loss : 1.6555218634868276
for n_estimators = 200 and max depth = 2
Log Loss : 1.4860554893539555
for n_estimators = 200 and max depth = 3
Log Loss : 1.3873613261741593
for n_estimators = 200 and max depth = 5
Log Loss : 1.348287034752093
for n_estimators = 200 and max depth = 10
Log Loss : 1.681743803335293
for n_estimators = 500 and max depth = 2

```

```

Log Loss : 1.5679719930684377
for n_estimators = 500 and max depth = 3
Log Loss : 1.457283610902369
for n_estimators = 500 and max depth = 5
Log Loss : 1.3692398051935435
for n_estimators = 500 and max depth = 10
Log Loss : 1.7320880746674154
for n_estimators = 1000 and max depth = 2
Log Loss : 1.5440446124172909
for n_estimators = 1000 and max depth = 3
Log Loss : 1.4515796481751302
for n_estimators = 1000 and max depth = 5
Log Loss : 1.3645504800040156
for n_estimators = 1000 and max depth = 10
Log Loss : 1.7175778617068542
For values of best alpha = 50 The train log loss is: 0.06593142888575093
For values of best alpha = 50 The cross validation log loss is:
1.2833788362382992
For values of best alpha = 50 The test log loss is: 1.38521174031375

```

```

[109]: clf = RandomForestClassifier(max_depth=max_depth[int(best_alpha%4)],
    ↳n_estimators=alpha[int(best_alpha/4)], criterion='gini',
    ↳max_features='auto', random_state=42)
predict_and_plot_confusion_matrix(train_x_responseCoding,
    ↳train_y, cv_x_responseCoding, cv_y, clf)

```

```

Log loss : 1.2833788362382996
Number of mis-classified points : 0.49624060150375937

```

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



```
[110]: clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)],
    ↳ criterion='gini', max_depth=max_depth[int(best_alpha%4)], random_state=42,
    ↳ n_jobs=-1)
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 1
no_feature = 27
predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index]).
    ↳ reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
```

```

print("Predicted Class Probabilities:", np.round(sig_clf.
↳predict_proba(test_x_responseCoding[test_point_index].reshape(1,-1)),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
for i in indices:
    if i<9:
        print("Gene is important feature")
    elif i<18:
        print("Variation is important feature")
    else:
        print("Text is important feature")

```

```

Predicted Class : 2
Predicted Class Probabilities: [[0.0133 0.455  0.1365 0.0202 0.028  0.0253
0.2711 0.0426 0.008 ]]
Actual Class : 7

```

```

-----
Variation is important feature
Variation is important feature
Variation is important feature
Gene is important feature
Variation is important feature
Variation is important feature
Variation is important feature
Text is important feature
Text is important feature
Text is important feature
Text is important feature
Text is important feature
Gene is important feature
Variation is important feature
Text is important feature
Gene is important feature
Gene is important feature
Gene is important feature
Variation is important feature
Text is important feature
Gene is important feature
Variation is important feature
Text is important feature
Text is important feature
Gene is important feature
Gene is important feature
Gene is important feature

```

5 Stacking model

```
[111]: clf1 = SGDClassifier(alpha=0.001, penalty='l2', loss='log',  
    ↪class_weight='balanced', random_state=0)  
clf1.fit(train_x_onehotCoding, train_y)  
sig_clf1 = CalibratedClassifierCV(clf1, method="sigmoid")  
  
clf2 = SGDClassifier(alpha=1, penalty='l2', loss='hinge',  
    ↪class_weight='balanced', random_state=0)  
clf2.fit(train_x_onehotCoding, train_y)  
sig_clf2 = CalibratedClassifierCV(clf2, method="sigmoid")  
  
clf3 = MultinomialNB(alpha=0.001)  
clf3.fit(train_x_onehotCoding, train_y)  
sig_clf3 = CalibratedClassifierCV(clf3, method="sigmoid")  
  
sig_clf1.fit(train_x_onehotCoding, train_y)  
print("Logistic Regression : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf1.  
    ↪predict_proba(cv_x_onehotCoding))))  
sig_clf2.fit(train_x_onehotCoding, train_y)  
print("Support vector machines : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf2.  
    ↪predict_proba(cv_x_onehotCoding))))  
sig_clf3.fit(train_x_onehotCoding, train_y)  
print("Naive Bayes : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf3.  
    ↪predict_proba(cv_x_onehotCoding))))  
print("-"*50)  
alpha = [0.0001,0.001,0.01,0.1,1,10]  
best_alpha = 999  
for i in alpha:  
    lr = LogisticRegression(C=i)  
    sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3],  
    ↪meta_classifier=lr, use_probabilities=True)  
    sclf.fit(train_x_onehotCoding, train_y)  
    print("Stacking Classifier : for the value of alpha: %f Log Loss: %0.3f" %  
    ↪(i, log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))))  
    log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))  
    if best_alpha > log_error:  
        best_alpha = log_error
```

Logistic Regression : Log Loss: 1.08

Support vector machines : Log Loss: 1.71

Naive Bayes : Log Loss: 1.27

Stacking Classifier : for the value of alpha: 0.000100 Log Loss: 1.818

Stacking Classifier : for the value of alpha: 0.001000 Log Loss: 1.719

Stacking Classifier : for the value of alpha: 0.010000 Log Loss: 1.306

Stacking Classifier : for the value of alpha: 0.100000 Log Loss: 1.116
Stacking Classifier : for the value of alpha: 1.000000 Log Loss: 1.316
Stacking Classifier : for the value of alpha: 10.000000 Log Loss: 1.579

```
[112]: lr = LogisticRegression(C=0.1)
sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3],
    ↪meta_classifier=lr, use_probas=True)
sclf.fit(train_x_onehotCoding, train_y)

log_error = log_loss(train_y, sclf.predict_proba(train_x_onehotCoding))
print("Log loss (train) on the stacking classifier :",log_error)


log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
print("Log loss (CV) on the stacking classifier :",log_error)

log_error = log_loss(test_y, sclf.predict_proba(test_x_onehotCoding))
print("Log loss (test) on the stacking classifier :",log_error)


print("Number of missclassified point :", np.count_nonzero((sclf.
    ↪predict(test_x_onehotCoding)- test_y))/test_y.shape[0])
plot_confusion_matrix(test_y=test_y, predict_y=sclf.
    ↪predict(test_x_onehotCoding))
```

Log loss (train) on the stacking classifier : 0.501541165064261
Log loss (CV) on the stacking classifier : 1.1160944198293306
Log loss (test) on the stacking classifier : 1.1955539699907047
Number of missclassified point : 0.362406015037594


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



6 Maximum voting Classifier

```
[113]: #Refer:http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.
        ↳VotingClassifier.html
from sklearn.ensemble import VotingClassifier
vclf = VotingClassifier(estimators=[('lr', sig_clf1), ('svc', sig_clf2), ('rf', sig_clf3)], voting='soft')
vclf.fit(train_x_onehotCoding, train_y)
print("Log loss (train) on the VotingClassifier :", log_loss(train_y, vclf.predict_proba(train_x_onehotCoding)))
print("Log loss (CV) on the VotingClassifier :", log_loss(cv_y, vclf.predict_proba(cv_x_onehotCoding)))
```

```

print("Log loss (test) on the VotingClassifier :", log_loss(test_y, vclf.
    ↳predict_proba(test_x_onehotCoding)))
print("Number of missclassified point :", np.count_nonzero((vclf.
    ↳predict(test_x_onehotCoding)- test_y))/test_y.shape[0])
plot_confusion_matrix(test_y=test_y, predict_y=vclf.
    ↳predict(test_x_onehotCoding))

```


Log loss (train) on the VotingClassifier : 0.8587983986980064

Log loss (CV) on the VotingClassifier : 1.1723541208907562

Log loss (test) on the VotingClassifier : 1.2405947426215598

Number of missclassified point : 0.3669172932330827


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



[]: